Lecture 19:

Heterogeneous Parallelism and Hardware Specialization

Parallel Computer Architecture and Programming
CMU 15-418/15-618, Spring 2017

Kanye West

Power
(My Beautiful Dark Twisted Fantasy)

"My songs address the most important architectural issues of the time."

- Kanye

I want to begin this lecture by reminding you…

That we observed in assignment 1 that a well-optimized parallel implementation of a compute-bound application is about 44 times faster on my quad-core laptop than the output of single-threaded C code compiled with gcc -03.


YINZER
PROCESSORS
You need to buy a
new computer...

You need to buy a computer system

Processor A

4 cores

Each core has sequential performance P

Processor B

16 cores

Each core has sequential performance P/2

All other components of the system are equal.

Which do you pick?

Rewrite Amdahl's law in terms of resource limits

$$\text{speedup}(f, n, r) = \frac{1}{\frac{1-f}{\text{perf}(r)} + \frac{f}{\text{perf}(r) \cdot \frac{n}{r}}}$$

Relative to processor with 1 unit of resources, n=1.
Assume perf(1) = 1

f = fraction of program that is parallelizable

n = total processing resources (e.g., transistors on a chip)

r = resources dedicated to each processing core,

(each of the n/r cores has sequential performance $perf(r)$

More general form of Amdahl's Law in terms of f, n, r

Two examples where n=16

$r_A = 4$

$r_B = 1$

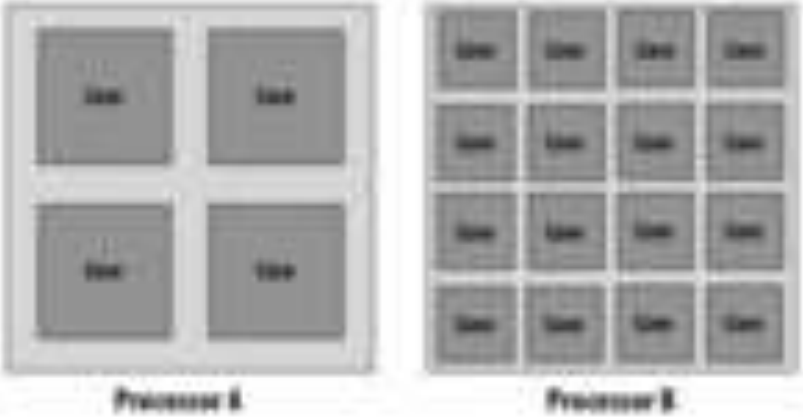

[Hill and Marty 08]

Speedup (relative to n=1)

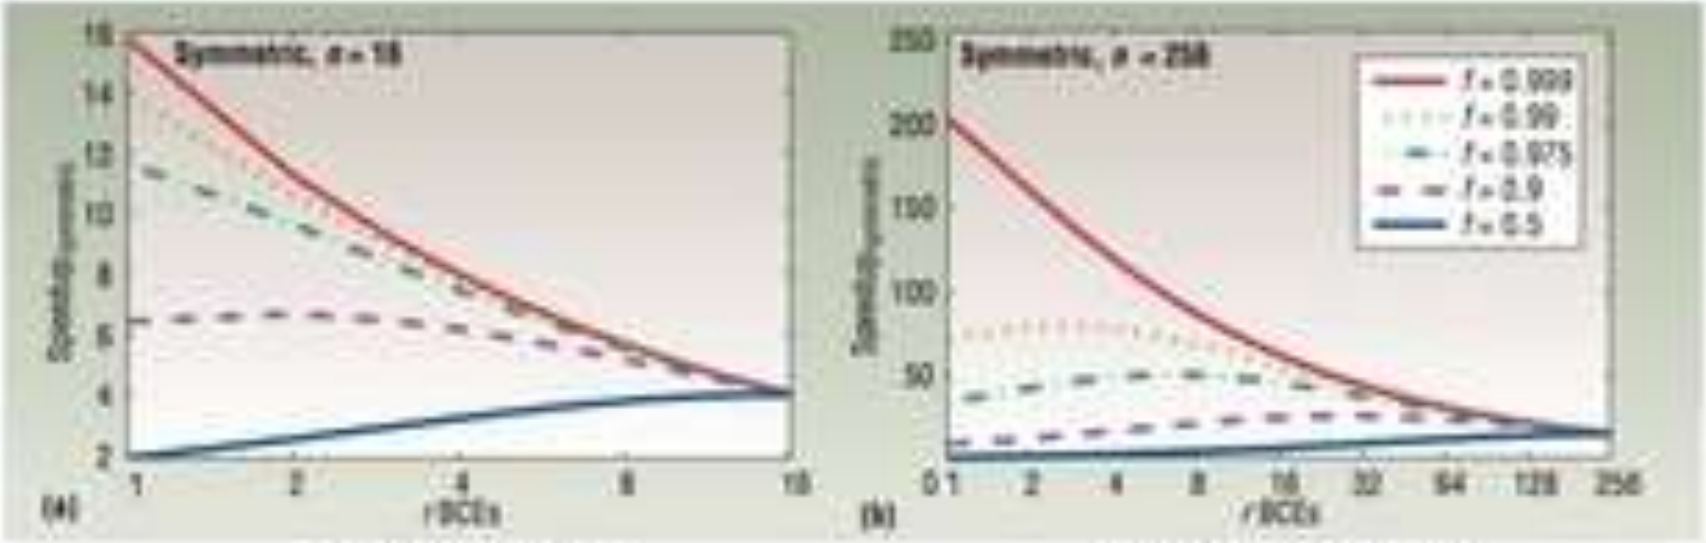

Up to 16 cores (n=16)

Up to 256 cores (n=256)

X-axis = r (chip with many small cores to left, fewer "fatter" cores to right)

Each line corresponds to a different workload

Each graph plots performance as resource allocation changes, but total chip resources resources kept the same (constant n per graph)

$perf(r)$ modeled as $\sqrt{r}$

[Figure credit: Hill and Marty 08]

Asymmetric set of processing cores

Example: n=16

One core: $r = 4$

Other 12 cores: $r = 1$

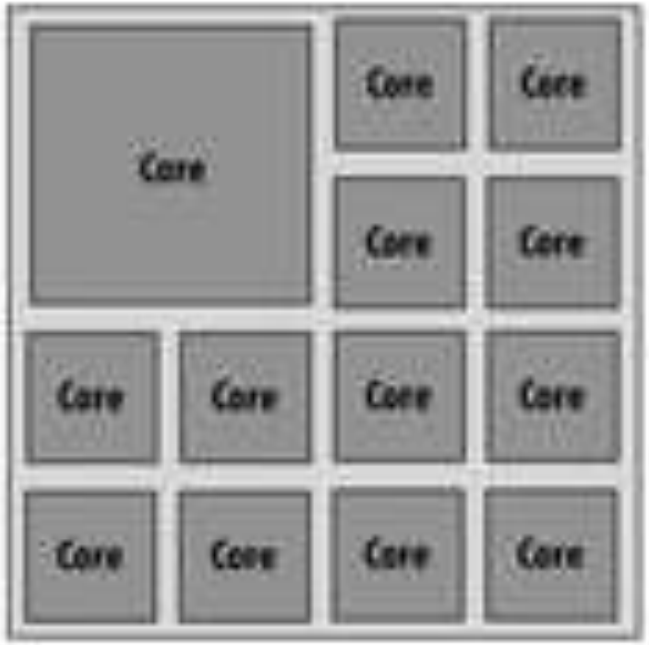

$$\text{speedup}(f, n, r) = \frac{1}{\frac{1-f}{\text{perf}(r)} + \underbrace{\frac{f}{\text{perf}(r)+(n-r)}}_{\text{one perf(r) processor + (n-r) perf(1)=1 processors}}}$$

(of heterogeneous processor with n recourses, relative to uniprocessor with one unit worth of resources, n=1)

[Hill and Marty 08]

Speedup (relative to n=1)

(Source: Hill and Marty 08)

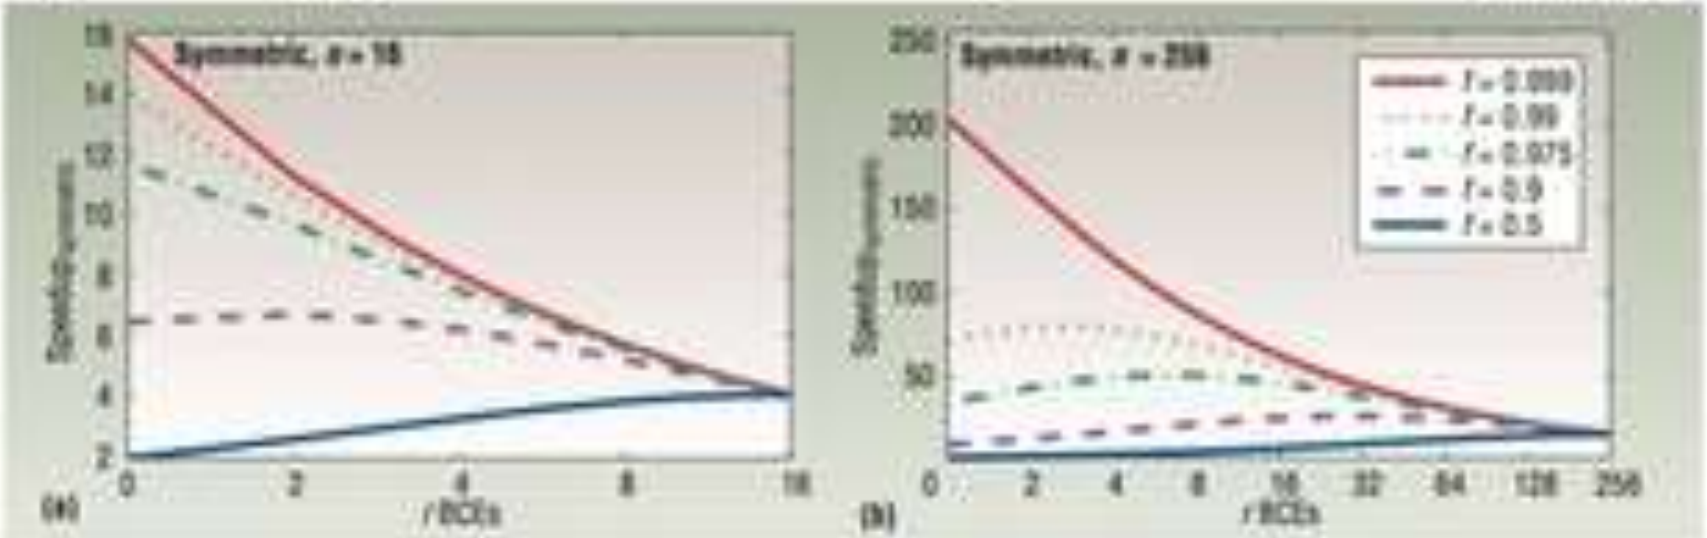

X-axis for symmetric architectures gives r for all cores (many small cores to left, few "fat" cores to right)

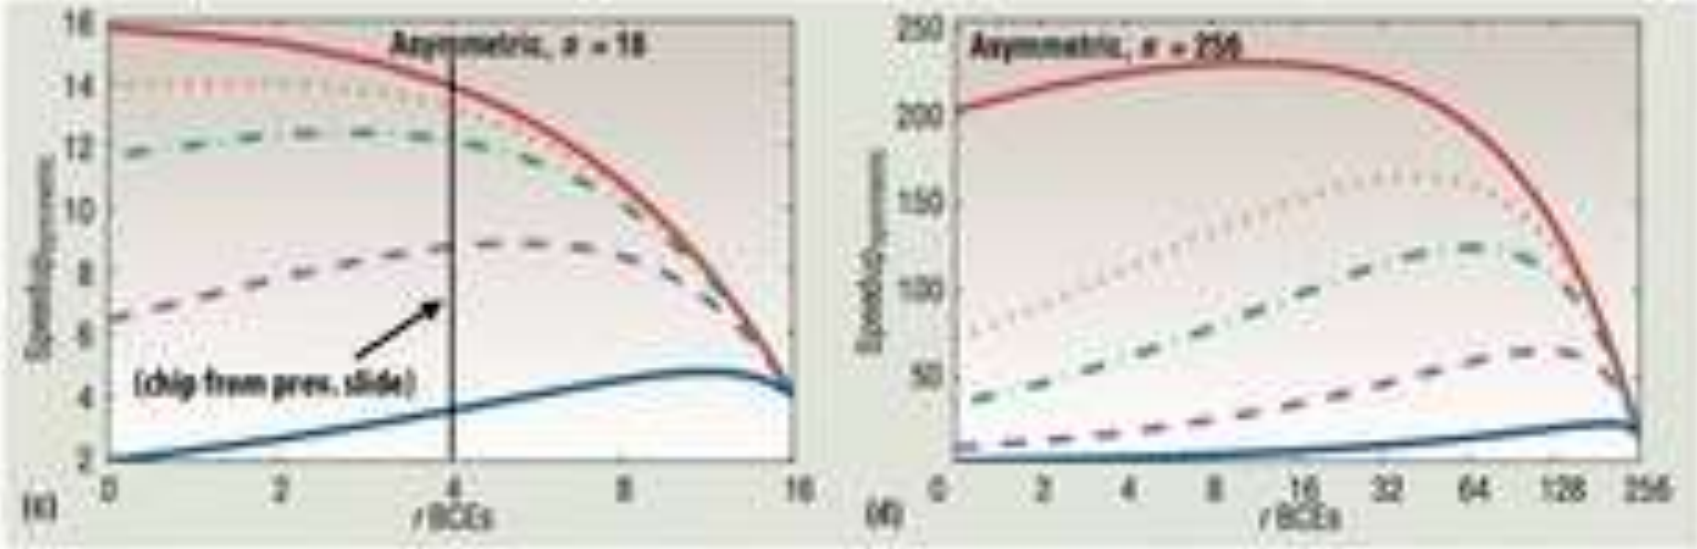

X-axis for asymmetric architectures gives r for the single "fat" core (assume rest of cores are $r = 1$)

Heterogeneous processing

Observation: most "real world" applications have complex workload characteristics *

They have components that can be widely parallelized.

And components that are difficult to parallelize.

They have components that are amenable to wide SIMD execution.

And components that are not. (divergent control flow)

They have components with predictable data access

And components with unpredictable access, but those accesses might cache well.

Idea: the most efficient processor is a heterogeneous mixture of resources ("use the most efficient tool for the job")

* You will likely make a similar observation during your projects

Example: Intel "Skylake" (2015)

(6th Generation Core i7 architecture)

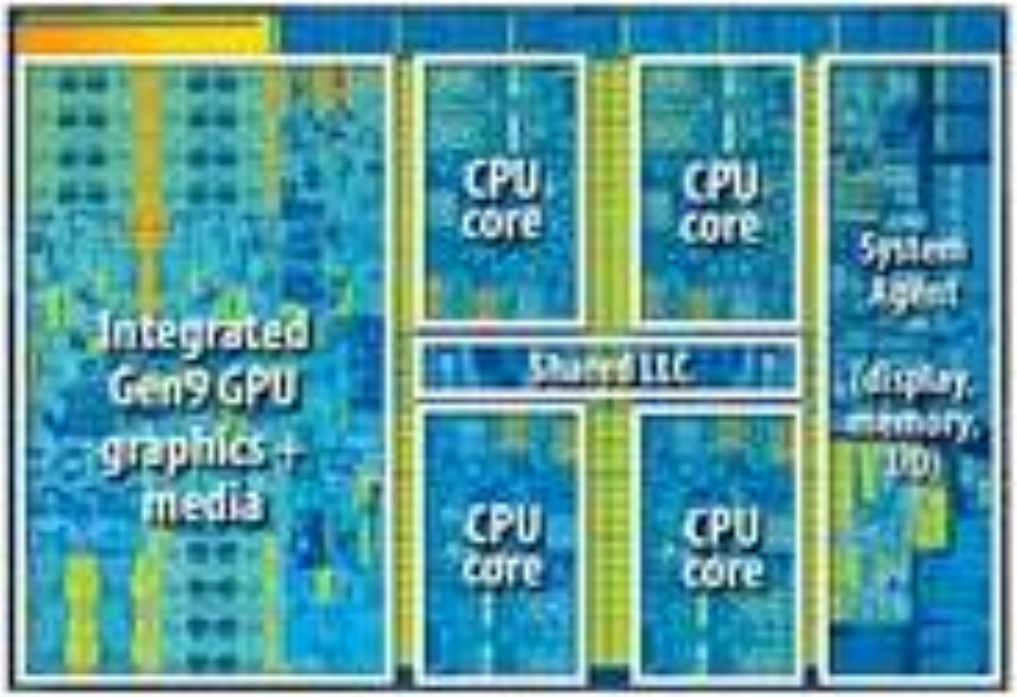

- CPU cores and graphics cores share same memory system
- Also share LLC (L3 cache)
  - Enables, low-latency, high-bandwidth communication between CPU and integrated GPU
- Graphics cores are cache coherent with CPU cores

More heterogeneity: add discrete GPU

Keep discrete (power hungry) GPU unless needed for graphics-intensive applications
Use integrated, low power graphics for basic graphics/window manager/UI

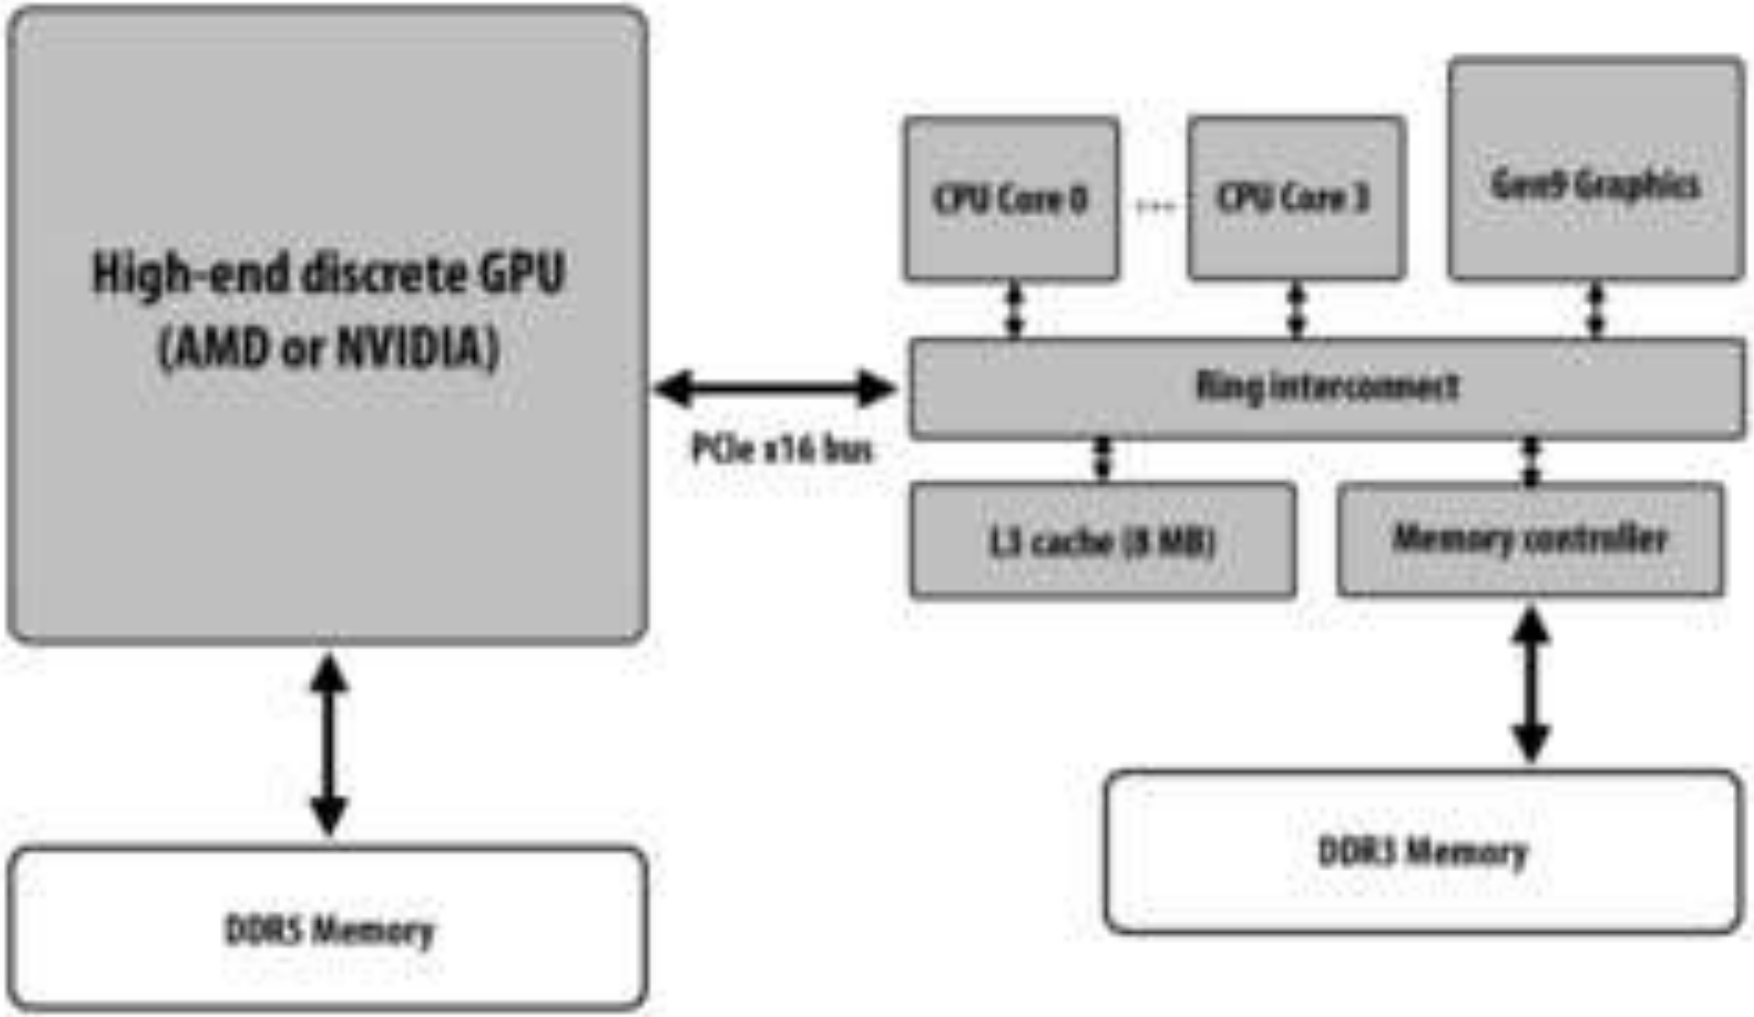

15in Macbook Pro 2011 (two GPUs)

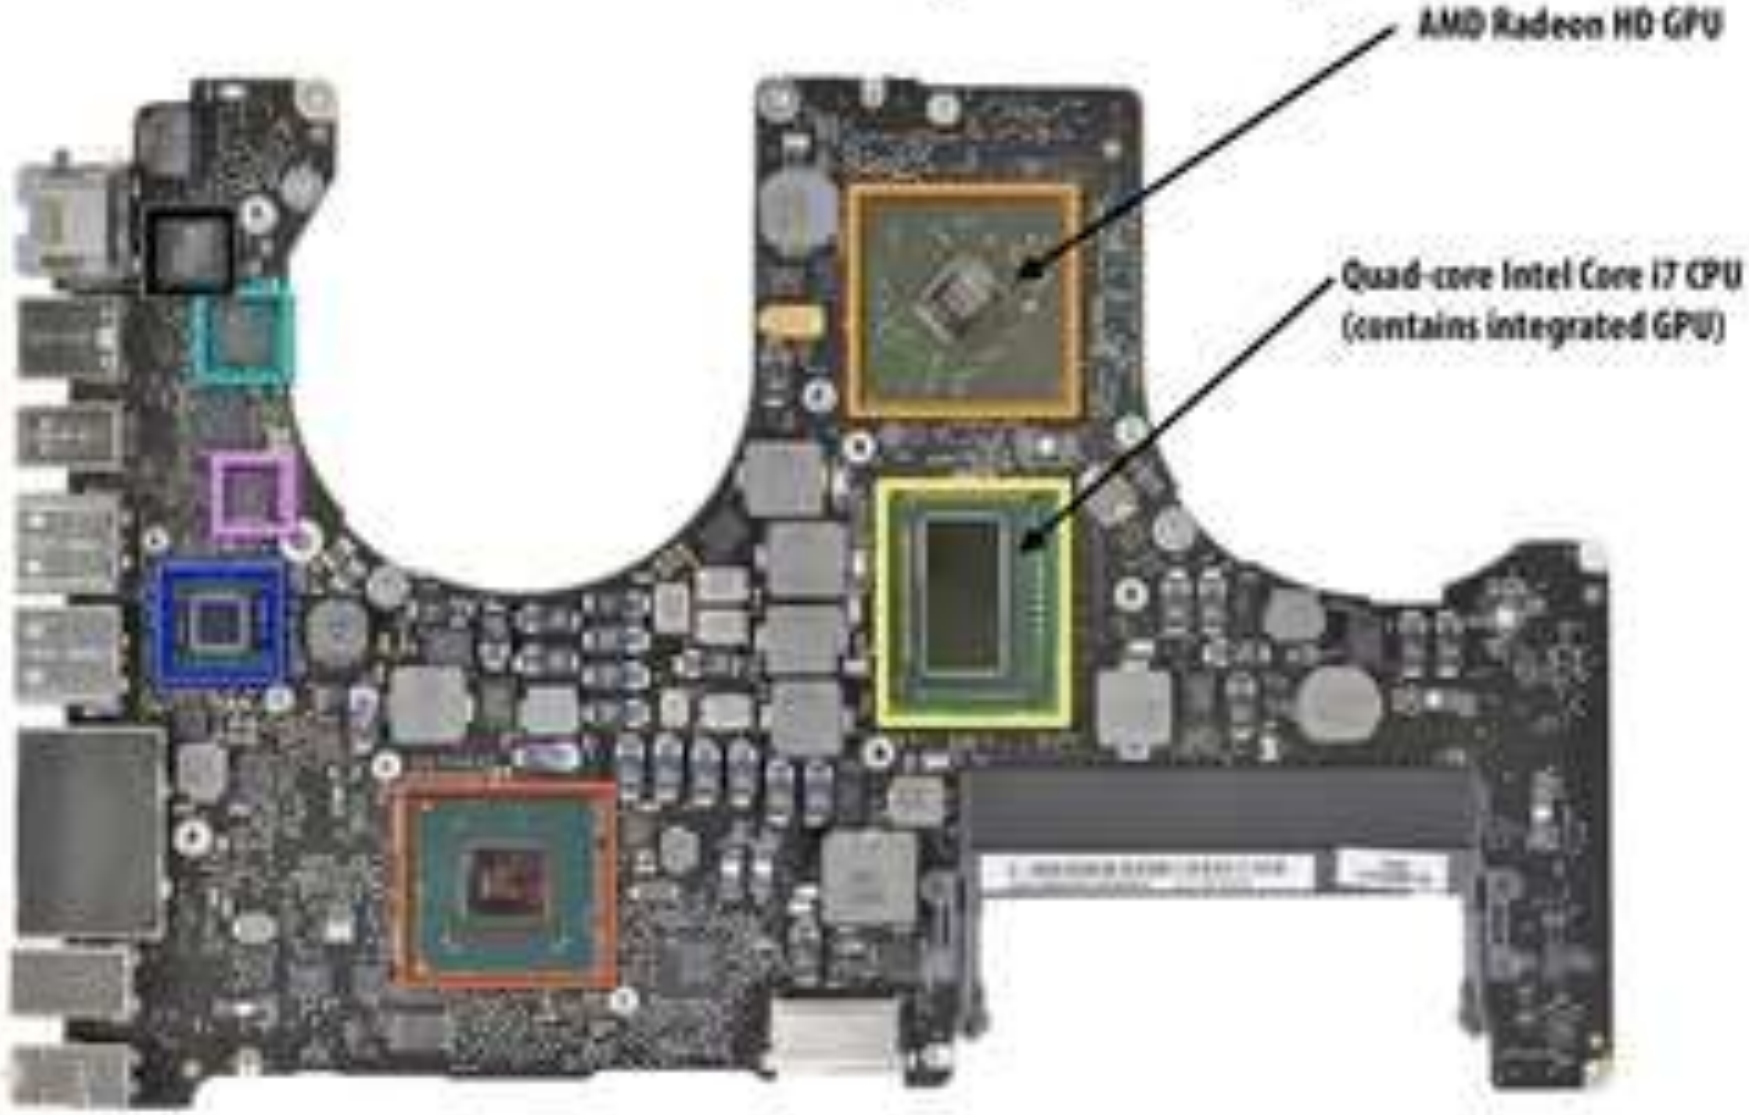

From ifixit.com teardown

Mobile heterogeneous processors

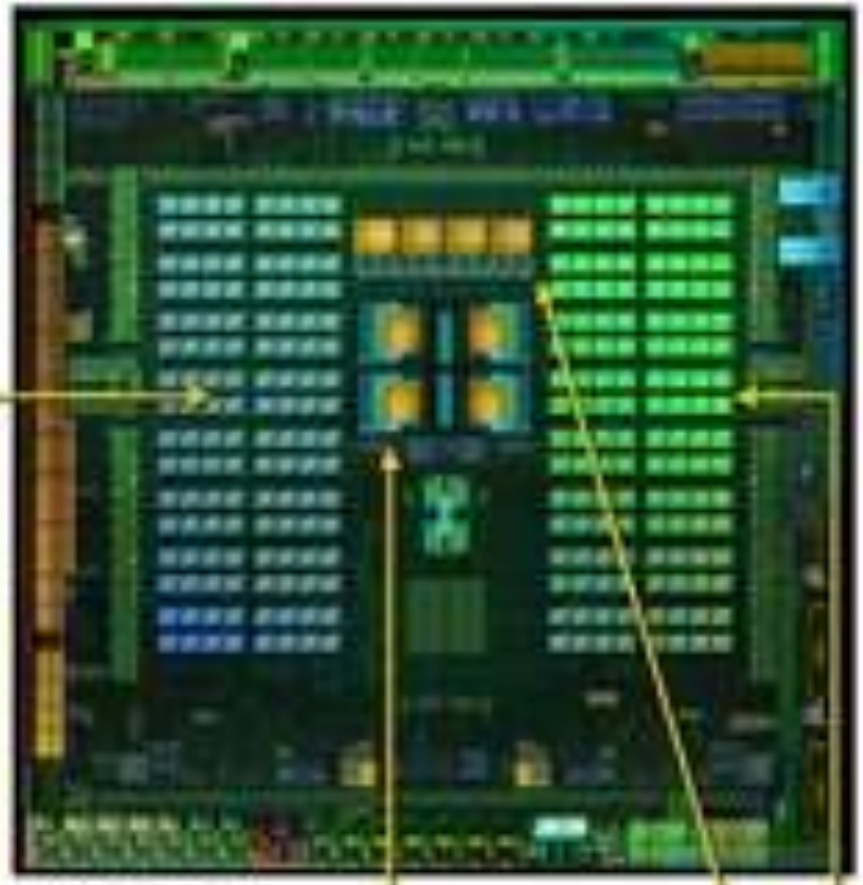

NVIDIA Tegra X1
Four ARM Cortex A57 CPU cores for applications
Four low performance (low power) ARM A53 CPU cores
One Maxwell SMM (256 "CUDA" cores)

Apple A9
Dual Core 64 bit CPU
GPU PowerVR GT6700 (6 "core") GPU

All image credit Chipworks, obtained from
http://www.anandtech.com/show/9686/the-apple-iphone-6s-and-iphone-6s-plus-review/3

Supercomputers use heterogeneous processing

Los Alamos National Laboratory: "Roadrunner"

Fastest US supercomputer in 2008, first to break Petaflop barrier: 1.7 PFLOPS

Unique at the time due to use of two types of processing elements

(IBM's Cell processor served as "accelerator" to achieve desired compute density)

- 6,480 AMD Opteron dual-core CPUs (12,960 cores)
- 12,970 IBM Cell Processors (1 CPU + 8 accelerator cores per Cell = 116,640 cores)
- 2.4 MWatt (about 2,400 average US homes)

GPU-accelerated supercomputing

- Oak Ridge Titan (world's #3)

- 18,688 AMD Opteron 16-core CPUs
- 18,688 NVIDIA Tesla K20X GPUs
- 710 TB RAM

- Estimated machine cost $97M
- Estimated annual power/operating cost: ~ $9M *

* Source: NPR

Intel Xeon Phi (Knights Landing)

- 72 "simple" x86 cores (1.1 Ghz, derived from Intel Atom)
- 16-wide vector instructions (AVX-512), four threads per core
- Targeted as an accelerator for supercomputing applications

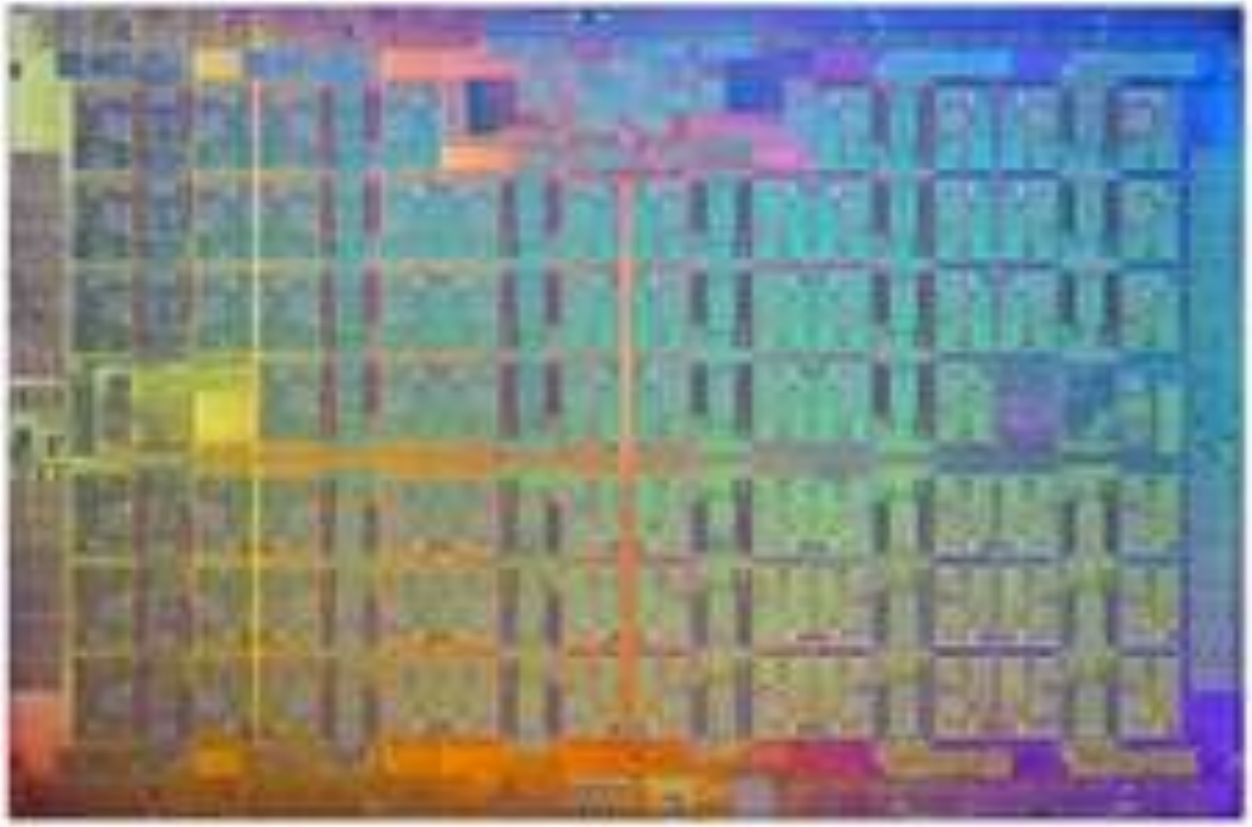

Heterogeneous architectures for supercomputing

Source: Top500.org Fall 2016 ranking

Rank	Site	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)	
1	National Supercomputing Center in Wuxi China	Sunway TaihuLight - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway NRCPC	10,649,600	93,014.6	125,435.9	15,371	93 PFLOPS, 15.3 MWatt (6078 MFLOPS/W)
2	National Super Computer Center in Guangzhou China	Tianhe-2 (MilkyWay-2) - TH-IVB-FEP Cluster, Intel Xeon E5-2692 12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P, NUDT	3,120,000	33,862.7	54,902.4	17,808	34 PFLOPS, 17.8 MWatt (1910 MFLOPS/W)
3	DOE/SC/Oak Ridge National Laboratory United States	Titan - Cray XK7 , Opteron 6274 16C 2.200GHz, Cray Gemini interconnect, NVIDIA K20x, Cray Inc.	560,640	17,590.0	27,112.5	8,209	18 PFLOPS, 8.2 MWatt (2143 MFLOPS/W)
4	DOE/NNSA/LLNL United States	Sequoia - BlueGene/Q, Power BQC 16C 1.60 GHz, Custom, IBM	1,572,864	17,173.2	20,132.7	7,890	
5	DOE/SC/LBNL/NERSC United States	Cori - Cray XC40, Intel Xeon Phi 7250 68C 1.4GHz, Aries interconnect, Cray Inc.	622,336	14,014.7	27,880.7	3,939	
6	Joint Center for Advanced High Performance Computing Japan	Oakforest-PACS - PRIMERGY CX1640 M1, Intel Xeon Phi 7250 68C 1.4GHz, Intel Omni-Path, Fujitsu	556,104	13,554.6	24,913.5	2,719	
7	RIKEN Advanced Institute for Computational Science (AICS) Japan	K computer, SPARC64 VIIIfx 2.0GHz, Tofu interconnect, Fujitsu	705,024	10,510.0	11,280.4	12,660	
8	Swiss National Supercomputing Centre (CSCS) Switzerland	Piz Daint - Cray XC30, Xeon E5-2670 8C 2.6GHz, Aries interconnect, NVIDIA Tesla K20X, Cray Inc.	206,720	9,779.0	15,988.0	1,312	

Xeon Phi

GPU

GPU

Green500: most energy efficient supercomputers

Efficiency metric: effective MFLOPS per Watt

Green500 Rank	MFLOPS/W	Site	System	Total Power(kW)
1	9462.1	NVIDIA Corporation	NVIDIA DGX-1, Xeon E5-2698v4 20C 2.2GHz, Infiniband EDR, NVIDIA Tesla P100	349.5
2	7453.5	Swiss National Supercomputing Centre (CSCS)	Cray XC50, Xeon E5-2690v3 12C 2.6GHz, Aries interconnect, NVIDIA Tesla P100	1312
3	6673.8	Advanced Center for Computing and Communication, RIKEN	ZettaScaler-1.6, Xeon E5-2618Lv3 8C 2.3GHz, Infiniband FDR, PEZY-SCnp	150.0
4	6051.3	National Supercomputing Center in Wuxi	Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway	15371
5	5806.3	Fujitsu Technology Solutions GmbH	PRIMERGY CX1640 M1, Intel Xeon Phi 7210 64C 1.3GHz, Intel Omni-Path	77
6	4985.7	Joint Center for Advanced High Performance Computing	PRIMERGY CX1640 M1, Intel Xeon Phi 7250 68C 1.4GHz, Intel Omni-Path	2718.7
7	4688.0	DOE/SC/Argonne National Laboratory	Cray XC40, Intel Xeon Phi 7230 64C 1.3GHz, Aries interconnect	1087
8	4112.1	Stanford Research Computing Center	Cray CS-Storm, Intel Xeon E5-2680v2 10C 2.8GHz, Infiniband FDR, Nvidia K80	190
9	4086.8	Academic Center for Computing and Media Studies (ACCMS), Kyoto University	Cray XC40, Intel Xeon Phi 7250 68C 1.4GHz, Aries interconnect	748.1
10	3836.4	Thomas Jefferson National Accelerator Facility	KDI Cluster, Intel Xeon Phi 7230 64C 1.3GHz, Intel Omni-Path	111

Source: Green500 Fall 2015 rankings

ARM-based supercomputers

- Observation: the heavy lifting in supercomputing applications is the data-parallel part of workload
  - Less need for "beefy" sequential performance cores

- Idea: build supercomputer out of power-efficient building blocks
  - ARM CPUs (for control/scheduling) + GPU cores or wide SIMD engines (serving as the primary compute engine)

- Montblanc: 64-bit ARM supercomputer (Barcelona Supercomputing Center)
  - www.montblanc-project.eu

- Fijitsu's new Petaflop-scale supercomputer (Post-K) based on ARMv8 (2020)

Also, although not a supercomputer, but Qualcomm announced an ARM-based server platform in December 2016 (48-core Centriq 2400)

Energy-constrained computing

- Supercomputers are energy constrained
  - Due to shear scale
  - Overall cost to operate (power for machine and for cooling)
- Datacenters are energy constrained
  - Reduce cost of cooling
  - Reduce physical space requirements
- Mobile devices are energy constrained
  - Limited battery life
  - Heat dissipation

Energy-constrained computing

Limits on chip power consumption

- General mobile processing rule: the longer a task runs the less power it can use
  - Processor's power consumption is limited by heat generated (efficiency is required for more than just maximizing battery life)

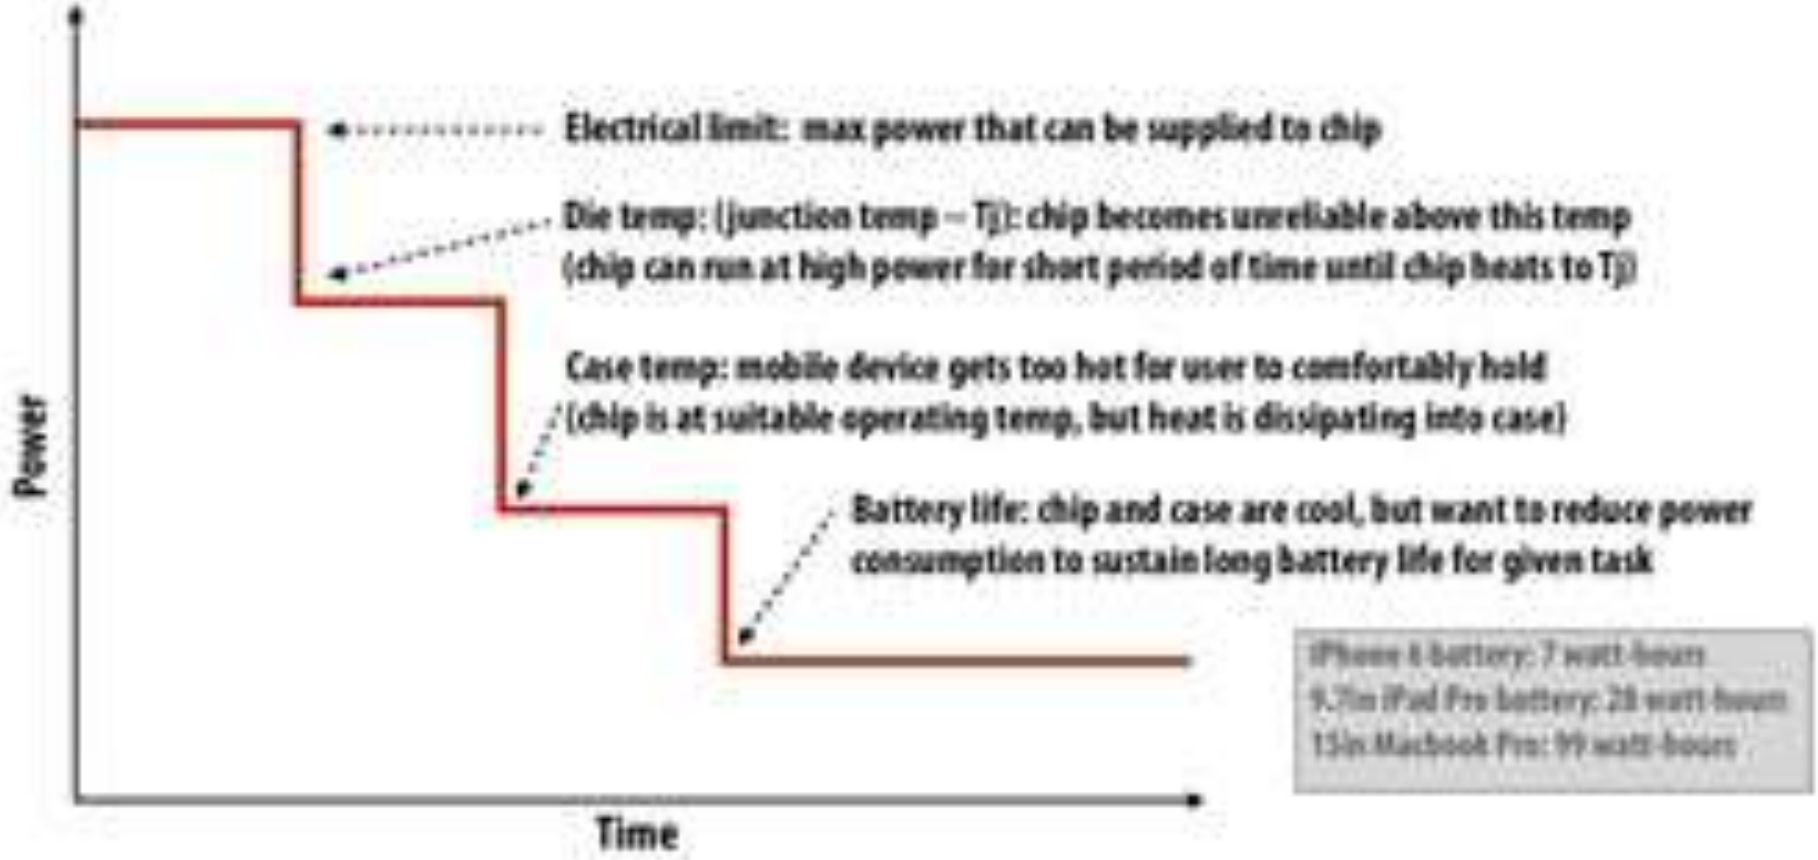

Slide credit: adopted from original slide from M. Shebanow: HPG 2013 keynote

Mobile: benefits of increasing efficiency

- Run faster for a fixed period of time
  - Run at higher clock, use more cores (reduce latency of critical task)
  - Do more at once
- Run at a fixed level of performance for longer
  - e.g., video playback, health apps
  - Achieve "always-on" functionality that was previously impossible

iPhone:
Siri activated by button press or holding phone up to ear

Amazon Echo / Google Home
Always listening

Google Glass: ~40 min recording per charge (nowhere near "always on")

Modern computing: efficiency often matters more than in the past, not less

Fourth, there's battery life.

To achieve long battery life when playing video, mobile devices must decode the video in hardware; decoding it in software uses too much power. Many of the chips used in modern mobile devices contain a decoder called H.264 – an industry standard that is used in every Blu-ray DVD player and has been adopted by Apple, Google (YouTube), Vimeo, Netflix and many other companies.

Although Flash has recently added support for H.264, the video on almost all Flash websites currently requires an older generation decoder that is not implemented in mobile chips and must be run in software. The difference is striking: on an iPhone, for example, H.264 videos play for up to 10 hours, while videos decoded in software play for less than 5 hours before the battery is fully drained.

When websites re-encode their videos using H.264, they can offer them without using Flash at all. They play perfectly in browsers like Apple's Safari and Google's Chrome without any plugins whatsoever, and look great on iPhones, iPods and iPads.

Steve Jobs' "Thoughts on Flash", 2010

http://www.apple.com/hotnews/thoughts-on-flash/

Pursuing highly efficient processing…
(specializing hardware beyond just parallel CPUs and GPUs)

Efficiency benefits of compute specialization

- Rules of thumb: compared to high-quality C code on CPU...
- Throughput-maximized processor architectures: e.g., GPU cores
  - Approximately 10x improvement in perf / watt
  - Assuming code maps well to wide data-parallel execution and is compute bound
- Fixed-function ASIC ("application-specific integrated circuit")
  - Can approach 100-1000x or greater improvement in perf/watt
  - Assuming code is compute bound and is not floating-point math

[Source: Chung et al. 2010, Dally 08]

Wait… this entire class we've been talking about making efficient use out of multi-core CPUs and GPUs…
and now you're telling me these platforms are "inefficient"?

Why is a "general-purpose processor" so inefficient?

Consider the complexity of executing an instruction on a modern processor…

Read instruction — Address translation, communicate with icache, access icache, etc.
Decode instruction — Translate op to uops, access uop cache, etc.
Check for dependencies/pipeline hazards
Identify available execution resource
Use decoded operands to control register file SRAM (retrieve data)
Move data from register file to selected execution resource
Perform arithmetic operation
Move data from execution resource to register file
Use decoded operands to control write to register file SRAM

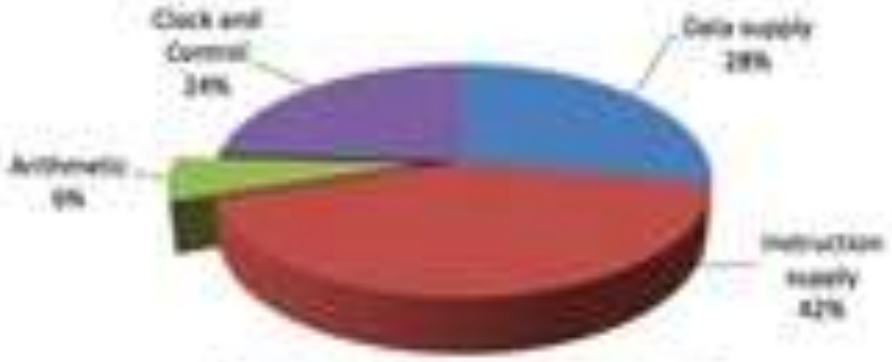

Review question:
How does SIMD execution reduce overhead of certain types of computations?
What properties must these computations have?

Efficient Embedded Computing (Daly et al. 08)
(Figure credit Eric Chung)

Contrast that complexity to the circuit required to actually perform the operation

Example: 8-bit logical OR

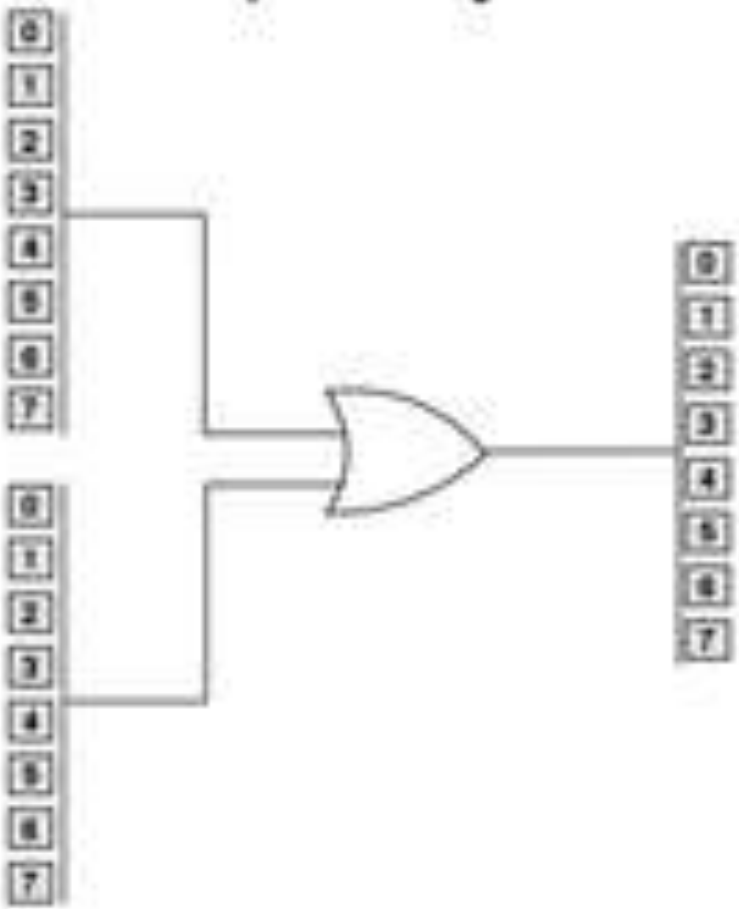

H.264 video encoding: fraction of energy consumed by different parts of instruction pipeline

[Hameed et al. ISCA 2010]

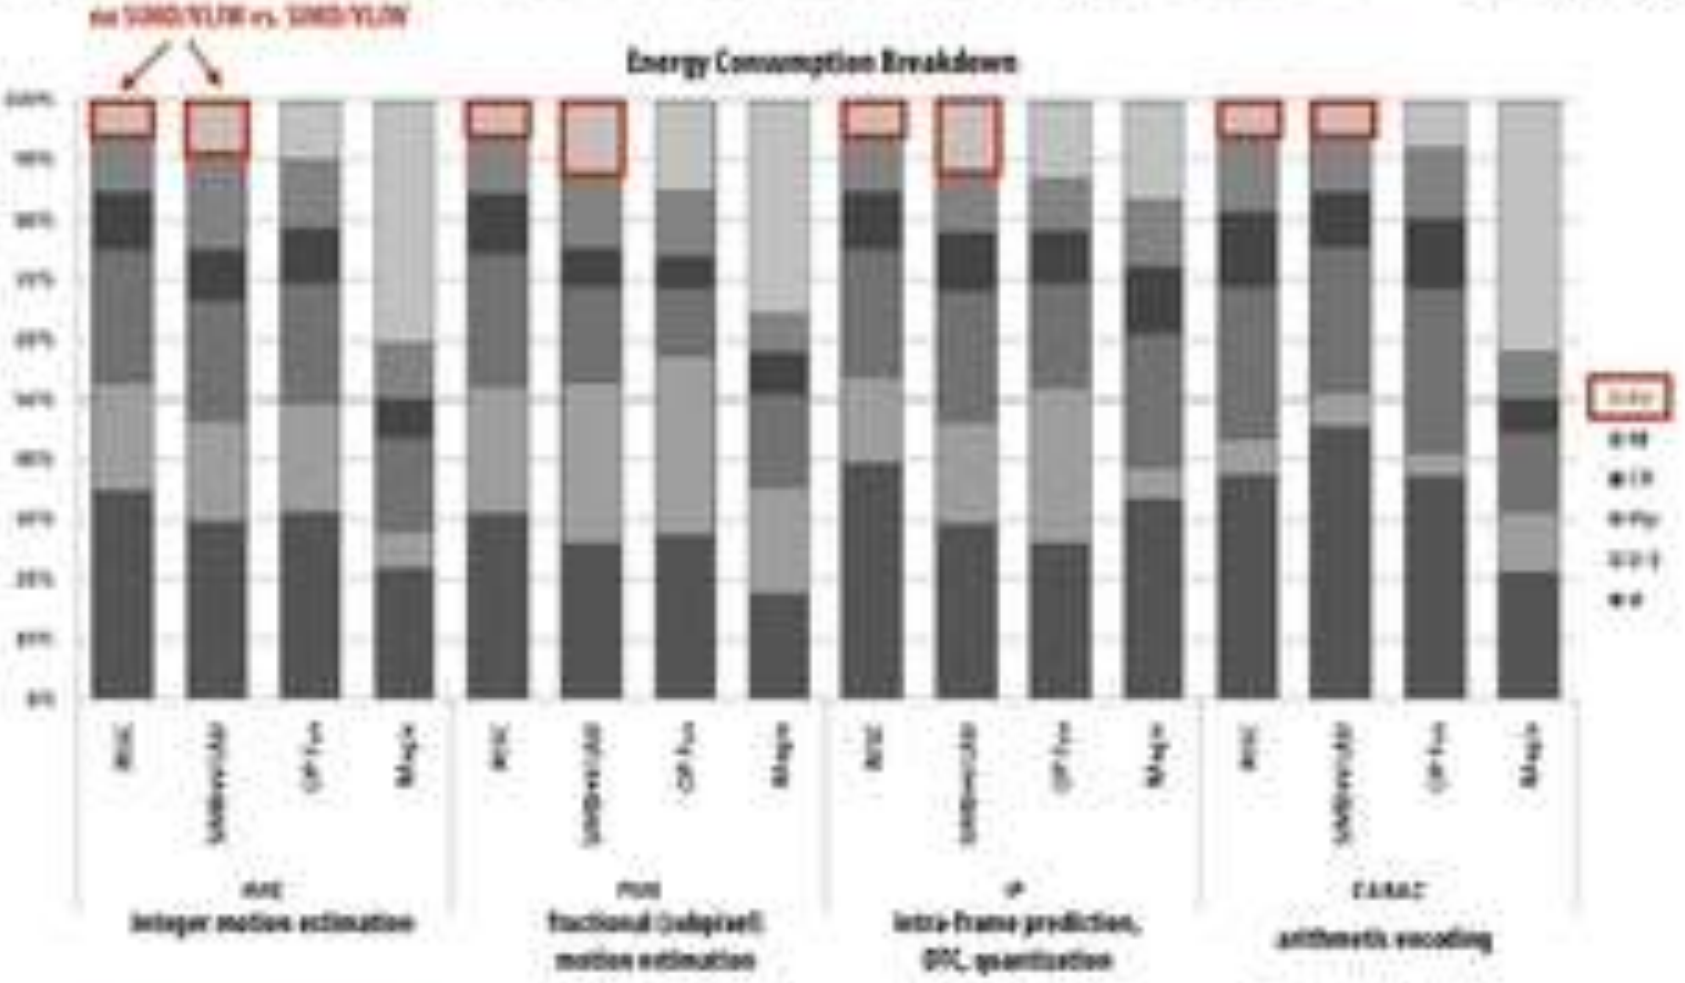

FFT: throughput/energy benefits of specialization

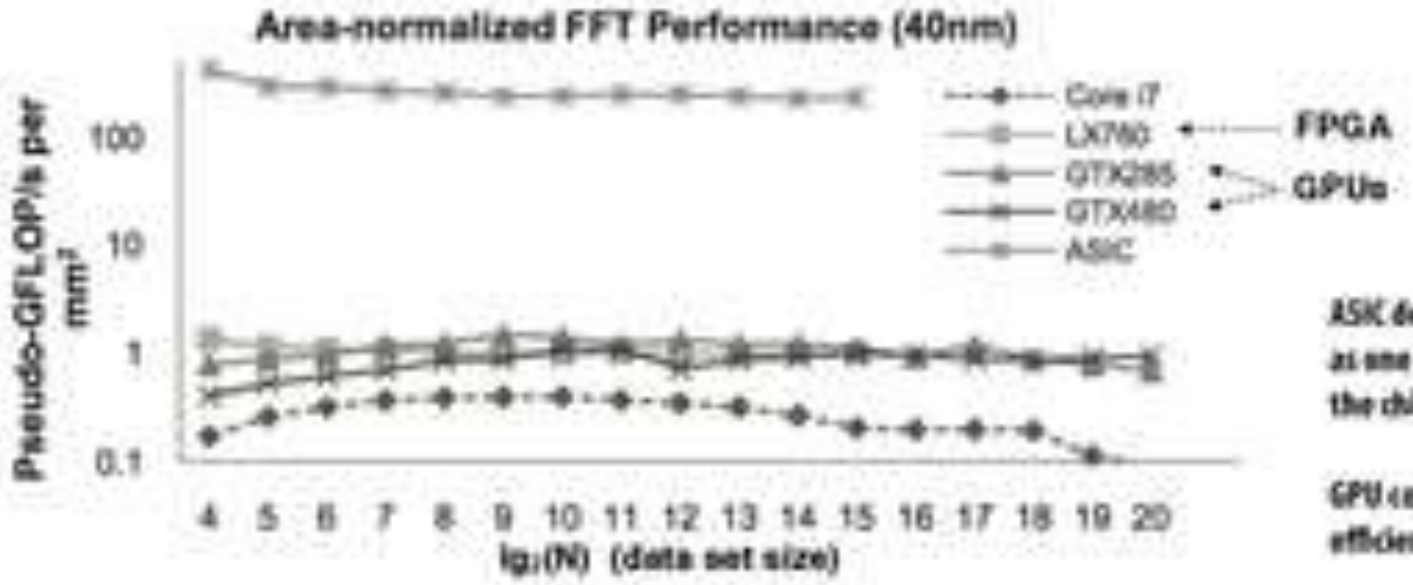

ASIC delivers same performance as one CPU core with ~ 1/1000th the chip area.

GPU cores: ~ 5-7 times more area efficient than CPU cores.

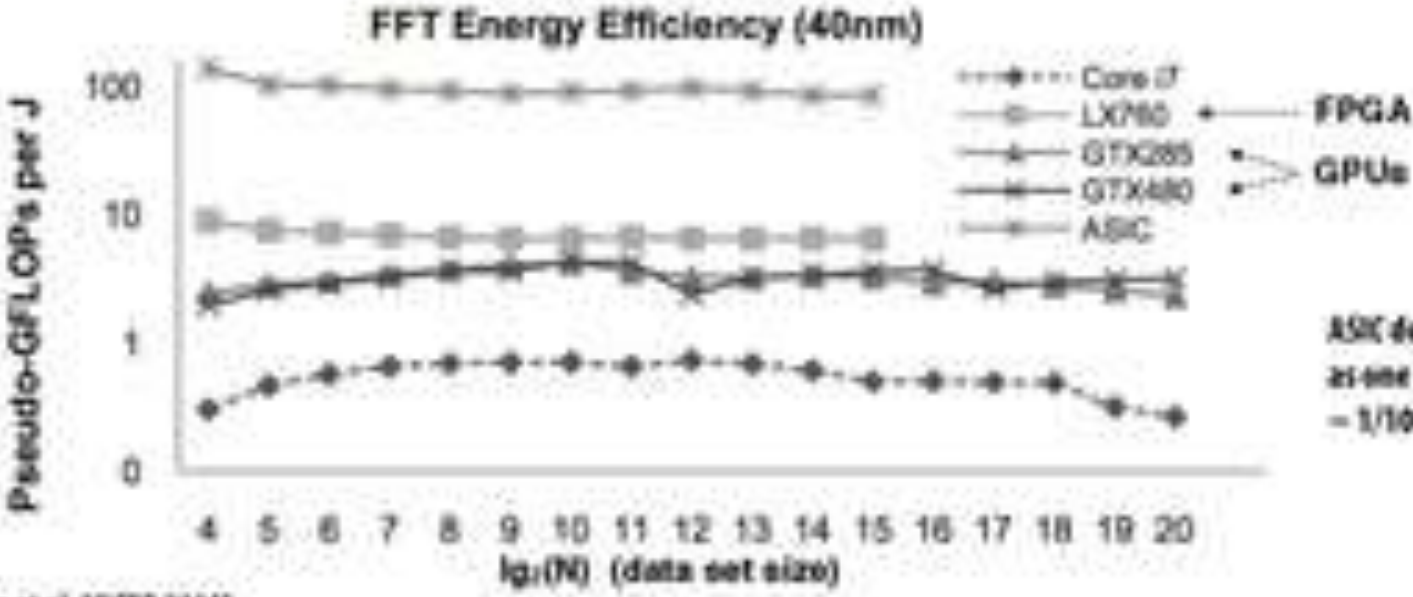

ASIC delivers same performance as one CPU core using only ~ 1/100th the power

[Chung et al. MICRO 2010]

GPU's are heterogeneous multi-core processors

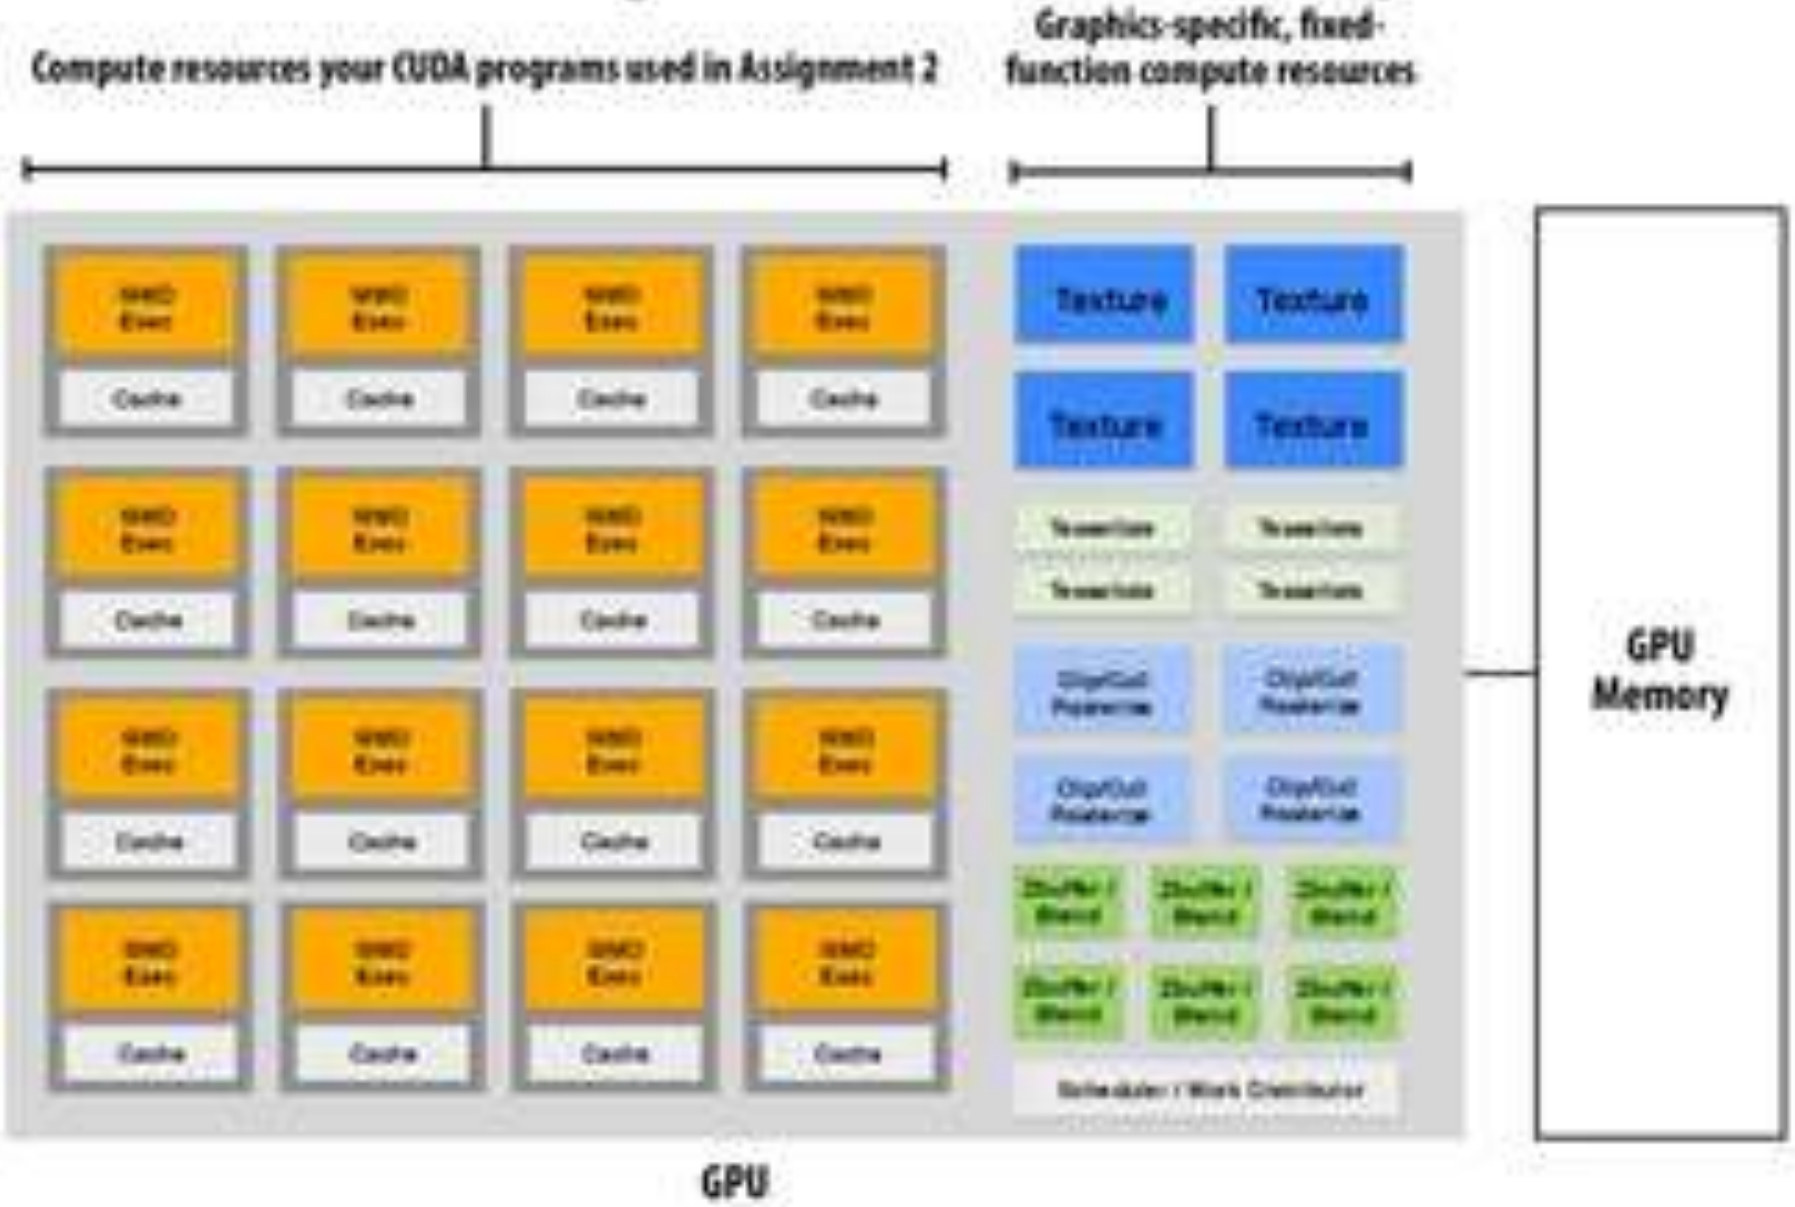

Example graphics tasks performed in fixed-function HW

Rasterization:
Determining what pixels a triangle overlaps

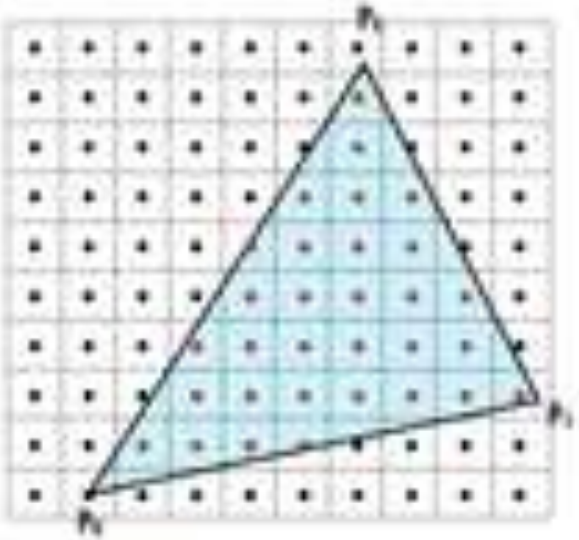

Texture mapping:
Warping/filtering images to apply detail to surfaces

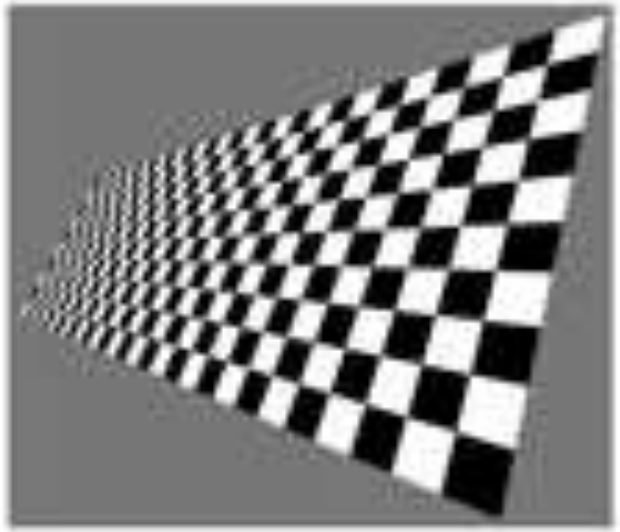

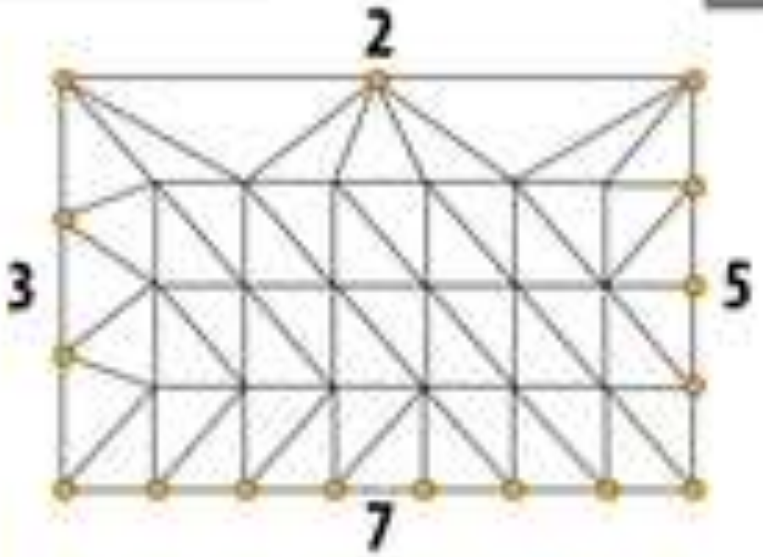

Geometric tessellation:
computing fine-scale geometry from coarse geometry

Anton supercomputer for molecular dynamics

(Developed by DE Shaw Research)

- Simulates time evolution of proteins
- ASIC for computing particle-particle interactions (512 of them in machine)
- Throughput-oriented subsystem for efficient fast-fourier transforms
- Custom, low-latency communication network designed for communication patterns of N-body simulations

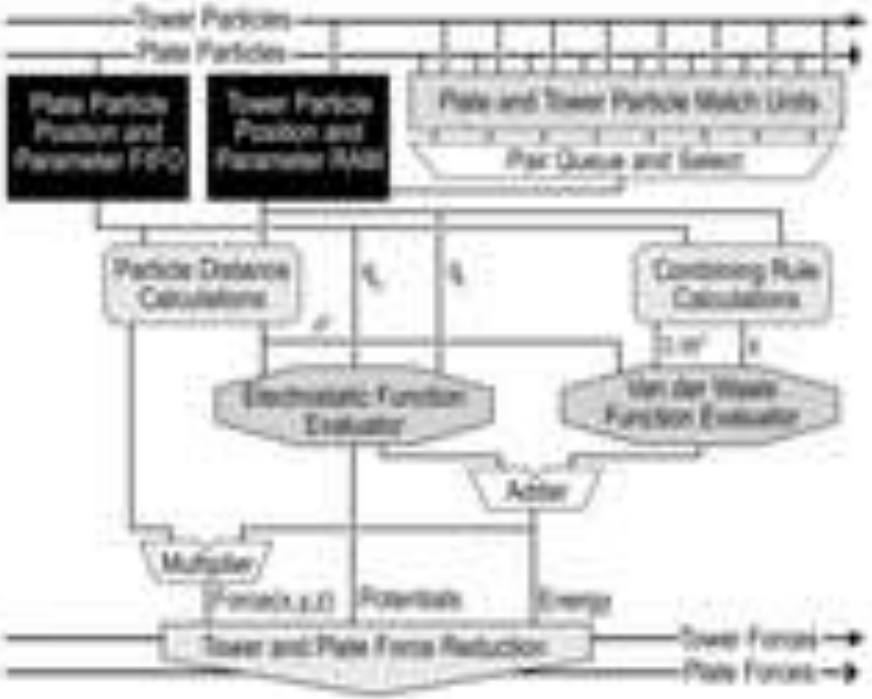

Specialized processors for evaluating deep networks

10+ papers at top computer architecture research conferences in 2016 on the topic of ASICs or accelerators for deep learning or evaluating deep networks…

Intel Lake Crest ML accelerator (formerly Nervana)

Digital signal processors (DSPs)

Programmable, but simpler instruction stream control paths

Complex instructions (e.g., SIMD/VLIW): perform many operations per instruction

Example: Qualcomm Hexagon DSP

Used for modem, audio, and (increasingly) image processing on Qualcomm Snapdragon SoC processors

VLIW: "very-long instruction word"
Single instruction specifies multiple different operations to do at once (contrast to SIMD)

Below: innermost loop of FFT
Hexagon DSP performs 29 "RISC" ops per cycle

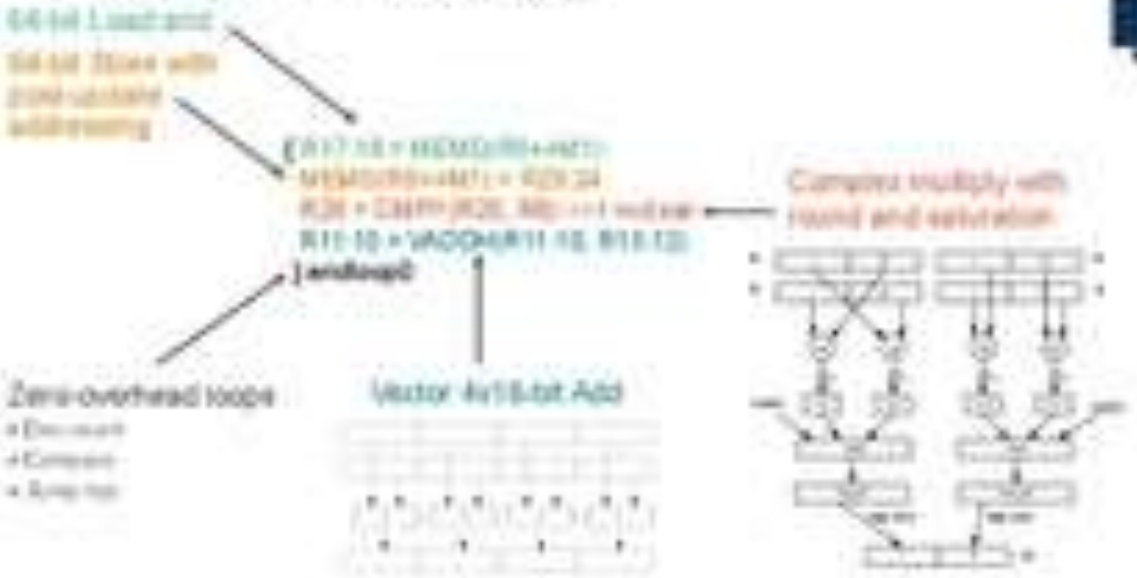

Hexagon DSP is in Google Pixel phone

Original iPhone touchscreen controller

Separate digital signal processor to interpret raw signal from capacitive touch sensor (do not burden main CPU)

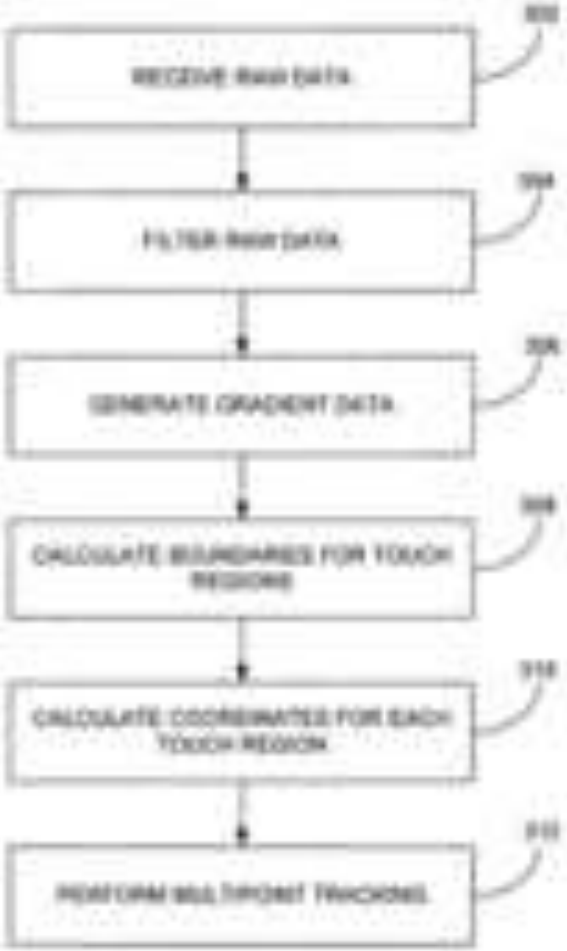

FIG. 16

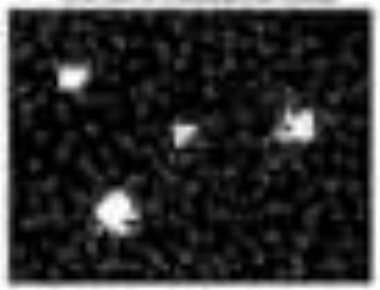

FIG. 17A

FIG. 17B

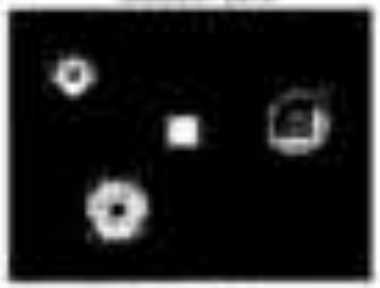

FIG. 17C

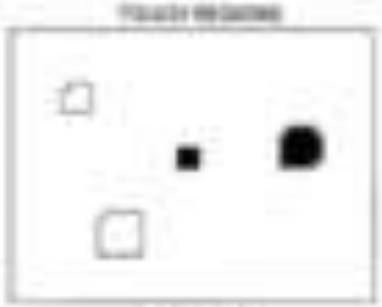

FIG. 17D

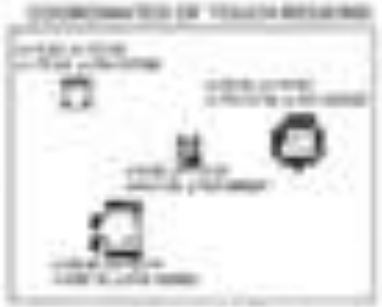

FIG. 17E

From US Patent Application 2006/0097991

Let's crack open a modern smartphone

Google Pixel Smartphone
Qualcomm Snapdragon 821 processor

Multi-core GPU
(3D graphics,
OpenCL data-parallel compute)

Display engine
(compresses pixels for
transfer to 4K screen)

Image Signal Processor (ISP)
ASIC for processing pixels off camera

Multi-core ARM CPU

Video encode/decode
ASIC (H.265 @ 4K)

"Hexagon"
Programmable DSP
data-parallel multi-media
processing

FPGAs (Field Programmable Gate Arrays)

- Middle ground between an ASIC and a processor
- FPGA chip provides array of logic blocks, connected by interconnect
- Programmer-defined logic implemented directly by FGPA

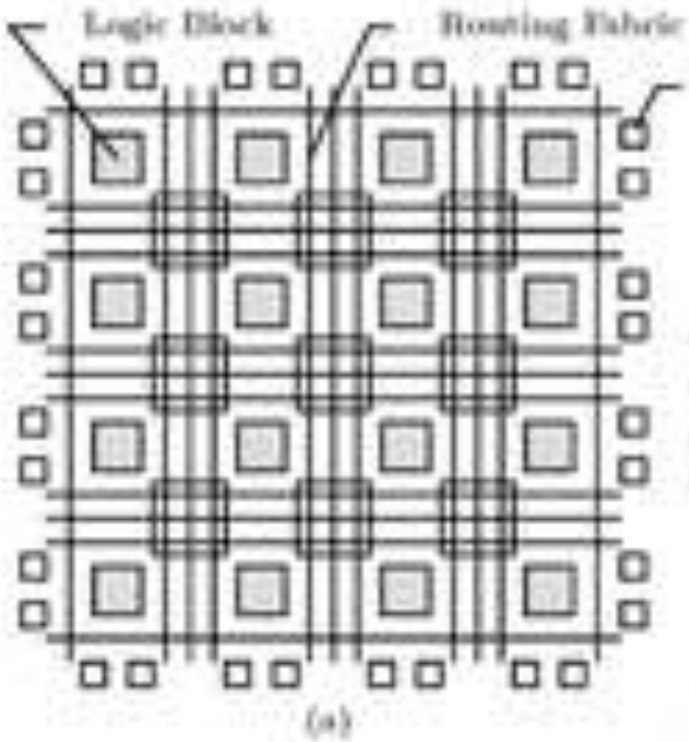

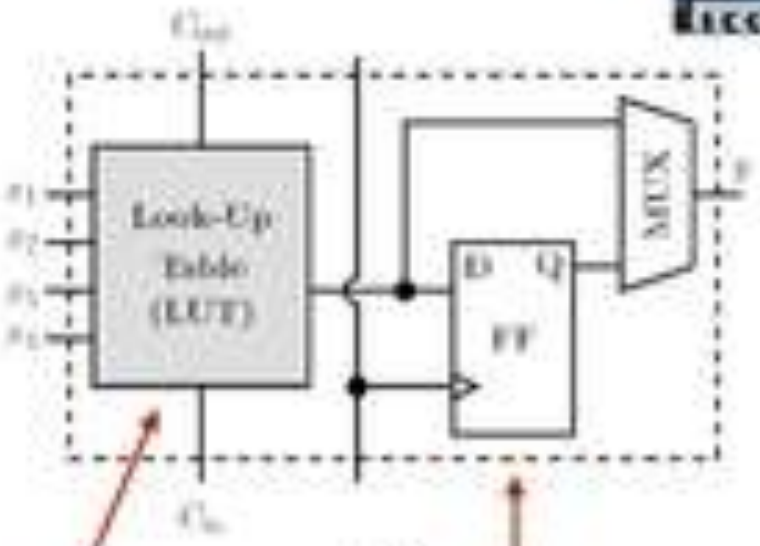

Flip flop (a register)

Programmable lookup table (LUT)

Image credit: Bai et al. 2014

Specifying combinatorial logic via LUT

- Example: 6-input, 1 output LUT in Xilinx Virtex-7 FPGAs
  - Think of a LUT6 as a 64 element table

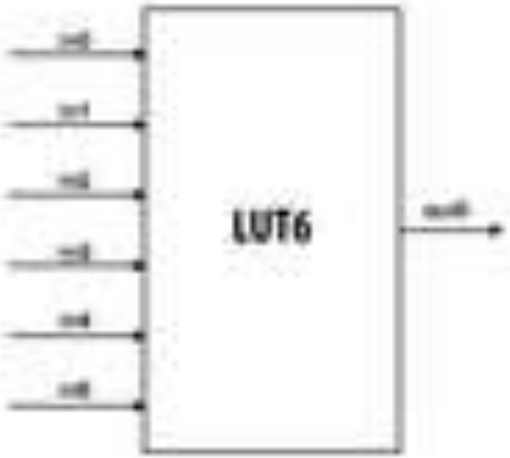

Example:
6-input AND

In	Out
0	0
1	0
2	0
3	0
⋮	⋮
63	1

40-input AND constructed by chaining outputs of eight LUT6's (delay = 3)

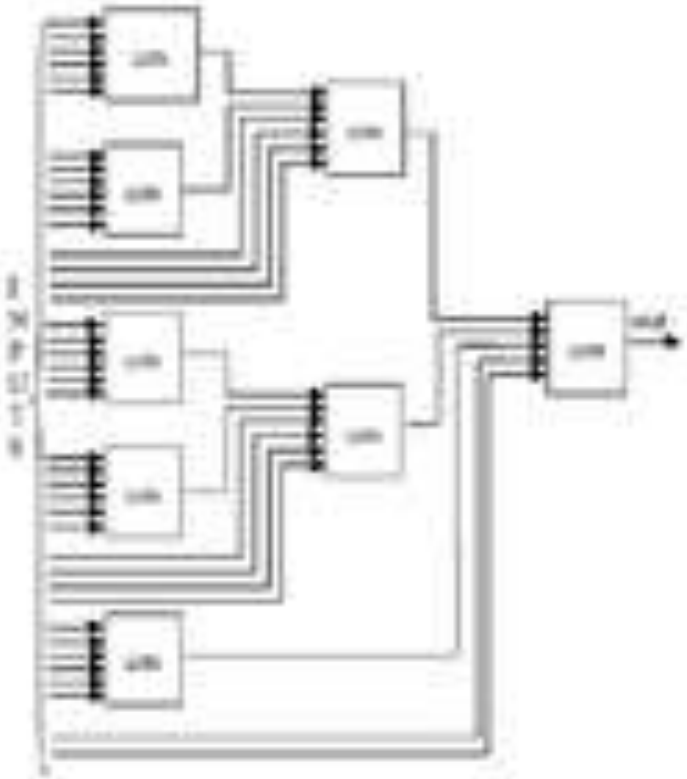

Image credit: [Zia 2011]

Project Catapult [Putnam et al. ISCA 2014]

- Microsoft Research investigation of use of FPGAs to accelerate datacenter workloads
- Demonstrated offload of part of Bing search's document ranking logic

FPGA board

1U server (Dual socket CPU + FPGA connected via PCIe bus)

Summary: choosing the right tool for the job

Energy-optimized CPU	Throughput-oriented processor (GPU)	Programmable DSP	FPGA/Future reconfigurable logic	ASIC
		HEXAGON		Video encode/decode, Audio playback, Camera RAW processing, neural nets (future?)
	~10X more efficient		~100X??? (jury still out)	~100-1000X more efficient
Easiest to program			Difficult to program (making it easier is active area of research)	Not programmable + costs 10-100's millions of dollars to design / verify / create

Credit Pat Hanrahan for this taxonomy

Challenges of heterogeneous designs:

(it's not easy to realize the potential of specialized, heterogeneous processing)

Challenges of heterogeneity

- Heterogeneous system: preferred processor for each task
- Challenge to software developer: how to map application onto a heterogeneous collection of resources?
  - Challenge: "Pick the right tool for the job": design algorithms that decompose into components that each map well to different processing components of the machine
  - The scheduling problem is more complex on a heterogeneous system
- Challenge for hardware designer: what is the right mixture of resources?
  - Too few throughput oriented resources (lower peak throughput for parallel workloads)
  - Too few sequential processing resources (limited by sequential part of workload)
  - How much chip area should be dedicated to a specific function, like video?

Pitfalls of heterogeneous designs

(Molnar 2010)

Consider a two stage graphics pipeline:
Stage 1: rasterize triangles into pixel fragments (using ASIC)
Stage 2: compute color of fragments (on SIMD cores)

Let's say you under-provision the rasterization unit on GPU:
Chose to dedicate 1% of chip area used for rasterizer to achieve throughput T fragments/clock
But really needed throughput of 1.2T to keep the cores busy (should have used 1.2% of chip area for rasterizer)

Now the programmable cores only run at 80% efficiency (99% of chip is idle 20% of the time = same perf as 79% smaller chip!)
So tendency is to be conservative and over-provision fixed-function components (diminishing their advantage)

Reducing energy consumption idea 1: use specialized processing

(use the right processor for the job)

Reducing energy consumption idea 2: move less data

Data movement has high energy cost

- Rule of thumb in mobile system design: always seek to reduce amount of data transferred from memory
  - Earlier in class we discussed minimizing communication to reduce stalls (poor performance). Now, we wish to reduce communication to reduce energy consumption
- "Ballpark" numbers [Sources: Bill Dally (NVIDIA), Tom Olson (ARM)]
  - Integer op: ~ 1 pJ *
  - Floating point op: ~20 pJ *
  - Reading 64 bits from small local SRAM (1mm away on chip): ~ 26 pJ
  - Reading 64 bits from low power mobile DRAM (LPDDR): ~1200 pJ ← Suggests that recomputing values, rather than storing and reloading them, is a better answer when optimizing code for energy efficiency!
- Implications
  - Reading 10 GB/sec from memory: ~1.6 watts
  - Entire power budget for mobile GPU: ~1 watt (remember phone is also running CPU, display, radios, etc.)
  - iPhone 6 battery: ~7 watt-hours (note: my Macbook Pro laptop: 99 watt-hour battery)
  - Exploiting locality matters!!!

* Cost to just perform the logical operation, not counting overhead of instruction decode, load data from registers, etc.

Three trends in energy-optimized computing

- **Compute less!**
  - Computing costs energy: parallel algorithms that do more work than sequential counterparts may not be desirable even if they run faster
- **Specialize compute units:**
  - Heterogeneous processors: CPU-like cores + throughput-optimized cores (GPU-like cores)
  - Fixed-function units: audio processing, "movement sensor processing" video decode/encode, image processing/computer vision?
  - Specialized instructions: expanding set of AVX vector instructions, new instructions for accelerating AES encryption (AES-NI)
  - Programmable soft logic: FPGAs
- **Reduce bandwidth requirements**
  - Exploit locality (restructure algorithms to reuse on-chip data as much as possible)
  - Aggressive use of compression: perform extra computation to compress application data before transferring to memory (likely to see fixed-function HW to reduce overhead of general data compression/decompression)

Summary: heterogeneous processing for efficiency

- Heterogeneous parallel processing: use a mixture of computing resources that fit mixture of needs of target applications
  - Latency-optimized sequential cores, throughput-optimized parallel cores, domain-specialized fixed-function processors
  - Examples exist throughout modern computing: mobile processors, servers, supercomputers
- Traditional rule of thumb in "good system design" is to design simple, general-purpose components
  - This is not the case in emerging systems (optimized for perf/watt)
  - Today: want collection of components that meet perf requirement AND minimize energy use
- Challenge of using these resources effectively is pushed up to the programmer
  - Current CS research challenge: how to write efficient, portable programs for emerging heterogeneous architectures?