

Lecture 23:

Efficiently Evaluating Deep Networks

Parallel Computer Architecture and Programming
CMU 15-418/15-618, Spring 2017

Something Good

Alt-J

(An Awesome Wave)

"Good things happen to students that come to class late in the semester."

- All musicians, everywhere.

Today

- High-performance evaluation of deep neural networks
- We will focus on the parallelism challenges of training deep networks next time

Training/evaluating deep neural networks

Technique leading to many high-profile AI advances in recent years

Speech recognition/natural language processing



Image interpretation and understanding




a baseball player swinging a bat at a ball
a boy is playing with a baseball bat



Game AI




Consider the following expression



$\max(\max(0, (a*b) + (c*d)) + (e*f) + (g*h), i*j)$

Consider expressing logic via a lookup table

- Example: 6-input, 1 output lookup table (LUT) in Xilinx Virtex-7 FPGAs
 - Think of a LUT6 as a 64 element table capable of expression functions of 6 bit inputs



Example:
6-input AND

in	Out
0	0
1	0
2	0
3	0
⋮	⋮
63	1

40-input AND constructed by chaining
outputs of eight LUT6's





Image credit: [24 2013]

What is a deep neural network?

A basic unit:

Unit with n inputs described by $n+1$ parameters
(weights + bias)



Example: rectified linear unit (ReLU)


$$f(x) = \max(0, x)$$

Basic computational interpretation:

It's just a circuit!

Biological inspiration:


unit output corresponds loosely to activation of neuron



Machine learning interpretation:

binary classifier: interpret output as the probability of one class

$$f(x) = \frac{1}{1 + e^{-x}}$$



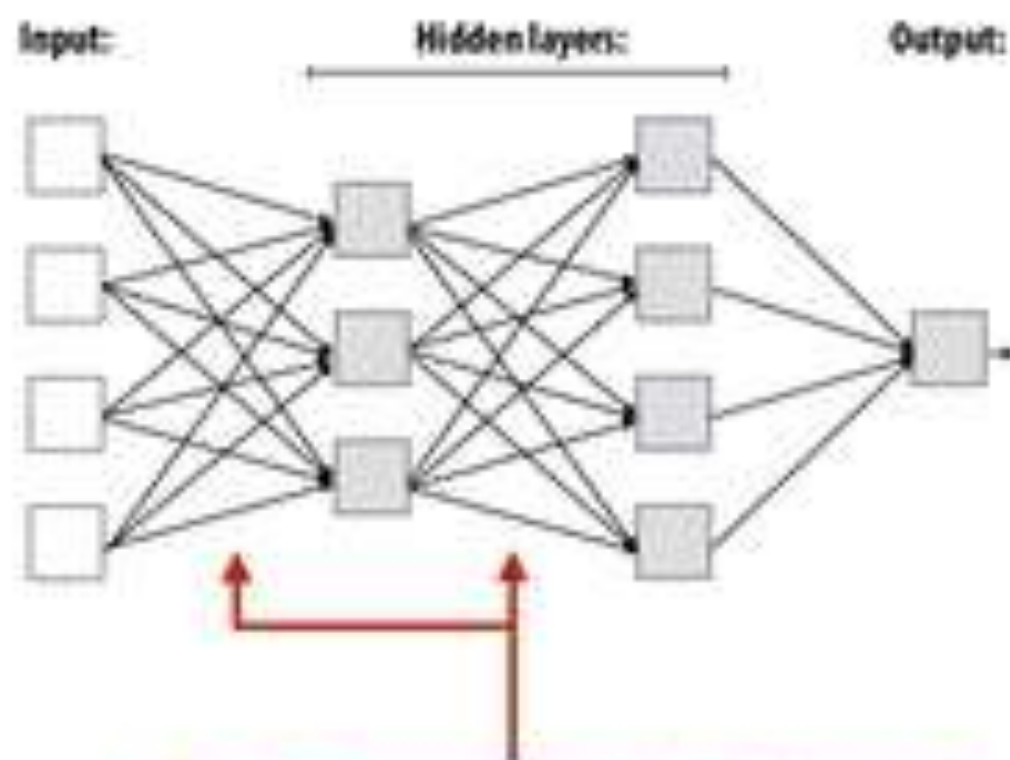
What is a deep neural network: topology

This network has: 4 inputs, 1 output, 7 hidden units

"Deep" = at least one hidden layer


Hidden layer 1: 3 units \times (4 weights + 1 bias) = 15 parameters

Hidden layer 2: 4 units \times (3 weights + 1 bias) = 16 parameters



Note "fully-connected" topology in this example
(every hidden unit receives input from all units on prior layer)

What is a deep neural network: topology




Recall image convolution (3x3 conv)

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1.0/9, 1.0/9, 1.0/9,
                  1.0/9, 1.0/9, 1.0/9,
                  1.0/9, 1.0/9, 1.0/9};
```

```
for (int j=0; j<HEIGHT; j++) {
  for (int i=0; i<WIDTH; i++) {
    float tmp = 0.f;
    for (int jj=0; jj<3; jj++)
      for (int ii=0; ii<3; ii++)
        tmp += input[(j+jj)*(WIDTH+2) + (i+ii)] * weights[jj*3 + ii];
    output[j*WIDTH + i] = tmp;
  }
}
```





Convolutional layer: locally connected AND all units in layer share the same parameters (same weights + same bias):
(note: network illustration only shows links for a 1D conv; a.k.a. one iteration of `ii` loop)

Strided 3x3 convolution

```
int WIDTH = 1024;
int HEIGHT = 1024;
int STRIDE = 2;
float input[(WIDTH+2) * (HEIGHT+2)];
float output[(WIDTH/STRIDE) * (HEIGHT/STRIDE)];

float weights[] = {1.0/9, 1.0/9, 1.0/9,
                   1.0/9, 1.0/9, 1.0/9,
                   1.0/9, 1.0/9, 1.0/9};

for (int j=0; j<HEIGHT; j+=STRIDE) {
  for (int i=0; i<WIDTH; i+=STRIDE) {
    float tmp = 0.f;
    for (int jj=0; jj<3; jj++)
      for (int ii=0; ii<3; ii++) {
        tmp += input[(j+jj)*(WIDTH+2) + (i+ii)] * weights[jj*3 + ii];
        output[(j/STRIDE)*WIDTH + (i/STRIDE)] = tmp;
      }
  }
}
```



What does convolution using these filter weights do?

$$\begin{bmatrix} .111 & .111 & .111 \\ .111 & .111 & .111 \\ .111 & .111 & .111 \end{bmatrix}$$

"Box blur"



What does convolution using these filter weights do?

$$\begin{bmatrix} .075 & .124 & .075 \\ .124 & .204 & .124 \\ .075 & .124 & .075 \end{bmatrix}$$

"Gaussian Blur"



What does convolution with these filters do?

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

Extracts horizontal
gradients

$$\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Extracts vertical
gradients

Gradient detection filters




Horizontal gradients



Vertical gradients

Note: you can think of a filter as a "detector" of a pattern, and the magnitude of a pixel in the output image as the "response" of the filter to the region surrounding each pixel in the input image

Applying many filters to an image at once




Applying many filters to an image at once


Input RGB image ($W \times H \times 3$)




96 (16×16) filters
(operate on RGB)



96 responses (normalized)




Adding additional layers



Example: "AlexNet" object detection network

Sequences of conv + reLU + pool (optional) layers

Example: AlexNet [Krizhevsky12]: 5 convolutional layers + 3 fully connected layers



Another example: VGG-16 [Simonyan15]: 13 convolutional layers

input: 224 x 224 RGB

conv/reLU: 3x3x3x64

conv/reLU: 3x3x64x64

maxpool

conv/reLU: 3x3x64x128

conv/reLU: 3x3x128x128

maxpool

conv/reLU: 3x3x128x256

conv/reLU: 3x3x256x256

conv/reLU: 3x3x256x256

maxpool

conv/reLU: 3x3x256x512

conv/reLU: 3x3x512x512

conv/reLU: 3x3x512x512

maxpool

conv/reLU: 3x3x512x512

conv/reLU: 3x3x512x512

conv/reLU: 3x3x512x512

maxpool

fully-connected-4096

fully-connected-4096

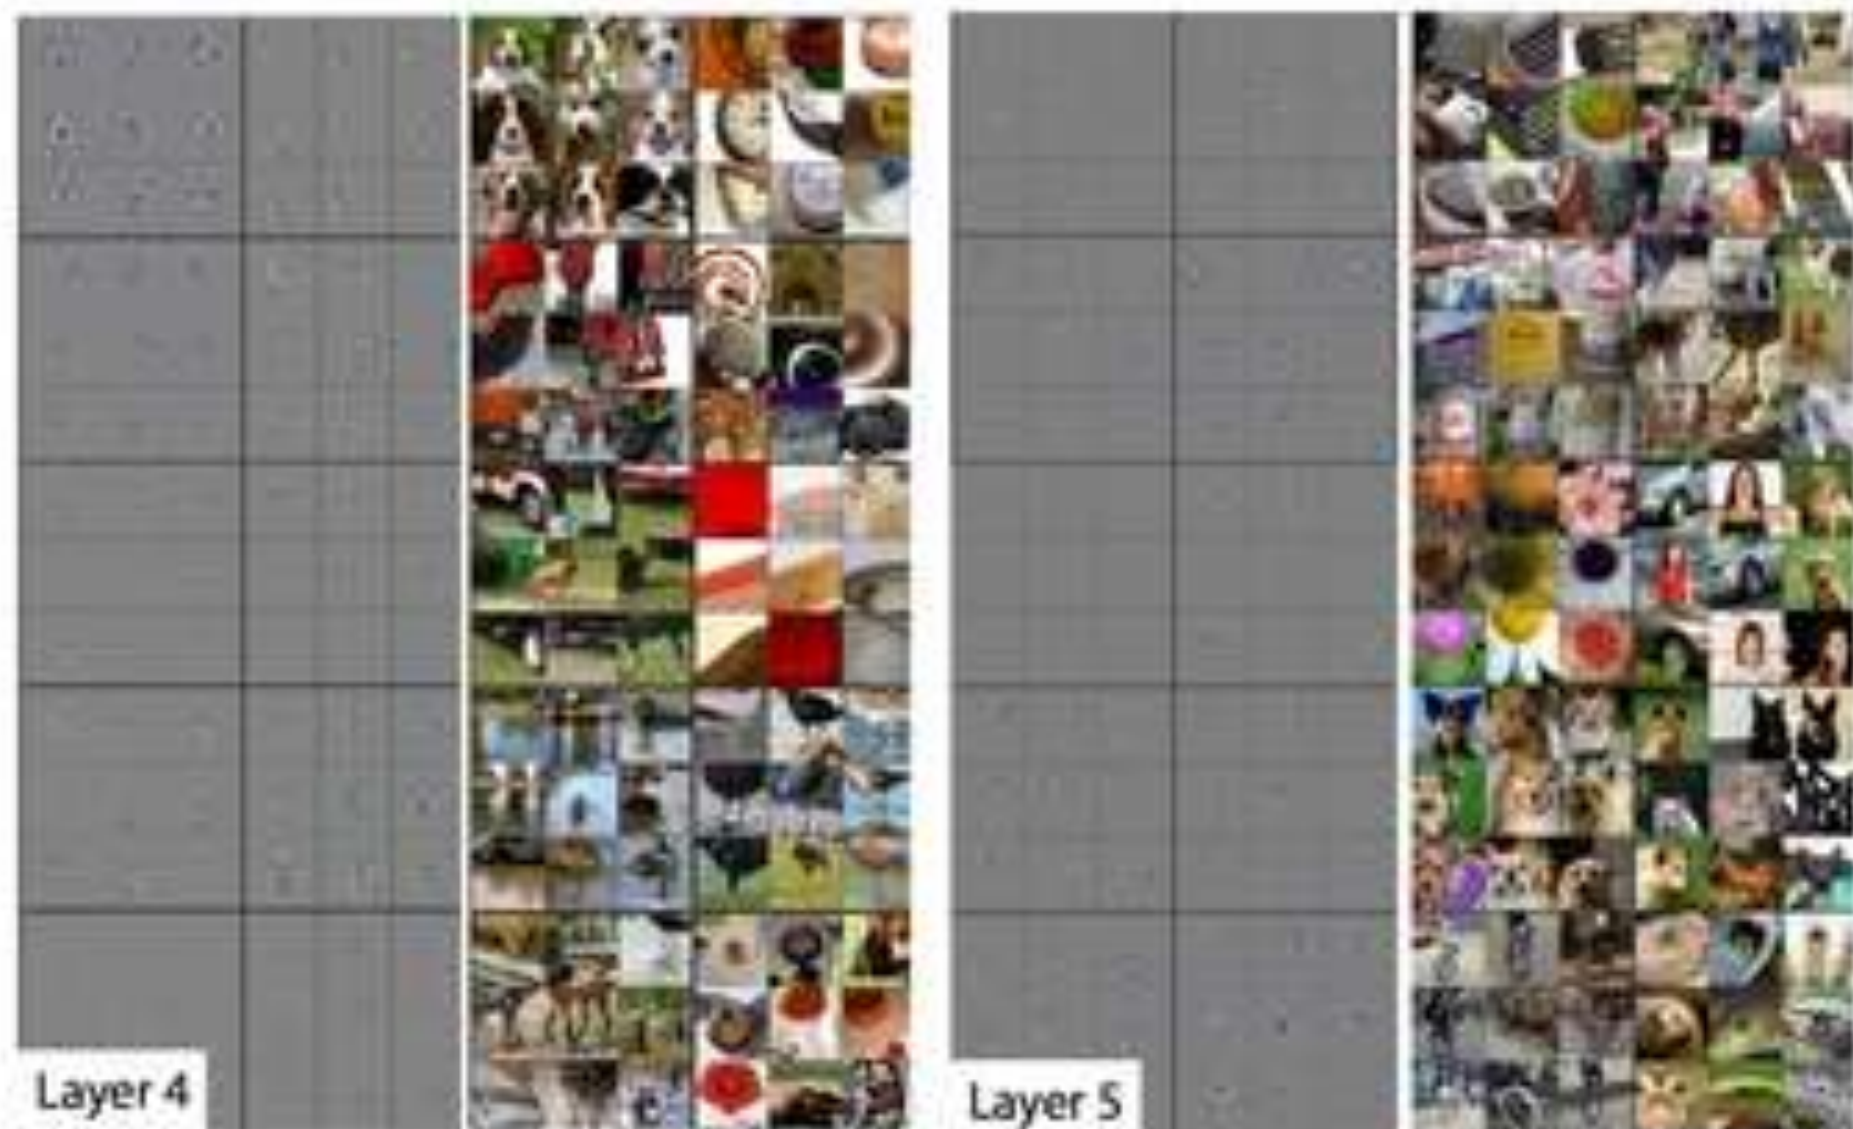
fully-connected-1000

soft-max

Why deep?




Why deep?




More recent image understanding networks



Inception (GoogleLeNet)



ResNet (34 layer version)



FCN for image segmentation

Deep networks learn useful representations


- Simultaneous, multi-scale learning of useful concepts for the task at hand
 - Example on previous slides: subparts detectors emerged in network for object classification
- But wait... how did you learn the values of all the weights?
 - Next class!
 - For today, assume the weights are given (today is about evaluating deep networks, not training them)

Efficiently implementing convolution layers

Dense matrix multiplication

```
float A[M][K];  
float B[K][N];  
float C[M][N];
```

```
// compute C += A * B  
#pragma omp parallel for  
for (int j=0; j<M; j++)  
    for (int i=0; i<N; i++)  
        for (int k=0; k<K; k++)  
            C[j][i] += A[j][k] * B[k][i];
```



What is the problem with this implementation?

Low arithmetic intensity (does not exploit temporal locality in access to A and B)

Blocked dense matrix multiplication

```
float A[M][K];  
float B[K][N];  
float C[M][N];
```

```
// compute C += A * B
```

```
#pragma omp parallel for
```

```
for (int jblock=0; jblock<M; jblock+=BLOCKSIZE_J)
```

```
    for (int iblock=0; iblock<N; iblock+=BLOCKSIZE_I)
```


```
        for (int kblock=0; kblock<K; kblock+=BLOCKSIZE_K)
```

```
            for (int j=0; j<BLOCKSIZE_J; j++)
```

```
                for (int i=0; i<BLOCKSIZE_I; i++)
```

```
                    for (int k=0; k<BLOCKSIZE_K; k++)
```

```
                        C[jblock+j][iblock+i] += A[jblock+j][kblock+k] * B[kblock+k][iblock+i];
```



Idea: compute partial result for block of C while required blocks of A and B remain in cache
(Assumes BLOCKSIZE chosen to allow block of A, B, and C to remain resident)

Self check: do you want as big a BLOCKSIZE as possible? Why?

Hierarchical blocked matrix mult

Exploit multi-level of memory hierarchy


```
float A[M][N];
float B[K][N];
float C[M][N];

// compute C += A * B
#pragma omp parallel for
for (int jblock2=0; jblock2<M; jblock2+=L2_BLOCKSIZE_J)
  for (int iblock2=0; iblock2<N; iblock2+=L2_BLOCKSIZE_I)
    for (int kblock2=0; kblock2<K; kblock2+=L2_BLOCKSIZE_K)
      for (int jblock1=0; jblock1<L1_BLOCKSIZE_J; jblock1+=L1_BLOCKSIZE_J)
        for (int iblock1=0; iblock1<L1_BLOCKSIZE_I; iblock1+=L1_BLOCKSIZE_I)
          for (int kblock1=0; kblock1<L1_BLOCKSIZE_K; kblock1+=L1_BLOCKSIZE_K)
            for (int j=0; j<BLOCKSIZE_J; j++)
              for (int i=0; i<BLOCKSIZE_I; i++)
                for (int k=0; k<BLOCKSIZE_K; k++)
                  ...
```

Not shown: final level of "blocking" for register locality...

Blocked dense matrix multiplication (1)

Consider SIMD parallelism
within a block




```
...
for (int j=0; j<BLOCKSIZE_J; j++) {
    for (int l=0; l<BLOCKSIZE_J; l+=SIMD_WIDTH) {
        simd_vec C_accum = vec_load(&C[jblock+j][lblock+l]);
        for (int k=0; k<BLOCKSIZE_K; k++) {
            // C = A*B + C
            simd_vec A_val = *pilat(&A[jblock+j][kblock+k]); // load a single element in vector register
            simd_madd(A_val, vec_load(&B[kblock+k][lblock+l]), C_accum);
        }
        vec_store(&C[jblock+j][lblock+l], C_accum);
    }
}
```

Vectorize i loop

Good: also improves spatial locality in access to B

Bad: working set increased by SIMD_WIDTH, still walking over B in large steps

Blocked dense matrix multiplication (2)



```

...
for (int j=0; j<BLOCKSIZE_J; j++)
    for (int i=0; i<BLOCKSIZE_J; i++) {
        float C_scalar = C[j*BLOCKSIZE_J+i][i*BLOCKSIZE_J];
        for (int k=0; k<BLOCKSIZE_K; k+=SIMD_WIDTH) {
            // C_scalar += dot(A,B)
            C_scalar += simd_dot(vect_load(A[j*BLOCKSIZE_J+i][k*BLOCKSIZE_J]), vect_load(BTtrans[i*BLOCKSIZE_J][k*BLOCKSIZE_J]));
        }
        C[j*BLOCKSIZE_J+i][i*BLOCKSIZE_J] = C_scalar;
    }
}


```

Assume i dimension is small. Previous vectorization scheme (1) would not work well.

Pre-transpose block of B (copy block of B to temp buffer in transposed form)

Vectorize innermost loop

Blocked dense matrix multiplication (3)



```
// assume blocks of A and B are row-transposed as Atrans and Btrans
for (int j=0; j<BLOCKSIZE_J; j+=SIMD_WIDTH) {
    for (int i=0; i<BLOCKSIZE_I; i+=SIMD_WIDTH) {

        simd_vec C_accum[SIMD_WIDTH];
        for (int k=0; k<SIMD_WIDTH; k++) // load C_accum for a SIMD_WIDTH x SIMD_WIDTH chunk of C^T
            C_accum[k] = vec_load(ACtrans[i*block+1+k][j*block+]);

        for (int k=0; k<BLOCKSIZE_K; k++) {
            simd_vec bvec = vec_load(BB[k*block+k][i*block+]);
            for (int kk=0; kk<SIMD_WIDTH; kk++) // innermost loop starts out dependent
                simd_madd(vec_load(BAtrans[k*block+k][j*block+]), splat(bvec[kk]), C_accum[kk]);
        }


        for (int k=0; k<SIMD_WIDTH; k++)
            vec_store(ACtrans[i*block+1+k][j*block+], C_accum[k]);
    }
}
```

Convolution as matrix-vector product

Construct matrix from elements of input image

x_{00}	x_{01}	x_{02}	x_{03}	...			
x_{10}	x_{11}	x_{12}	x_{13}	...			
x_{20}	x_{21}	x_{22}	x_{23}	...			
x_{30}	x_{31}	x_{32}	x_{33}	...			
...				

$O(N)$ storage multiplier for filter with N elements
Must construct input data matrix



Note: 0-pad matrix

3x3 convolution as matrix-vector product

Construct matrix from elements of input image


[illegible]

$O(N)$ storage overhead for filter with N elements
Must construct input data matrix

The diagram illustrates a 2D array structure with dimensions $W \times H$ and 9 columns. The array is divided into rows of different colors (blue, green, yellow, white, purple) and is associated with a vector of labels $[U_0, U_1, \dots, U_8]$.

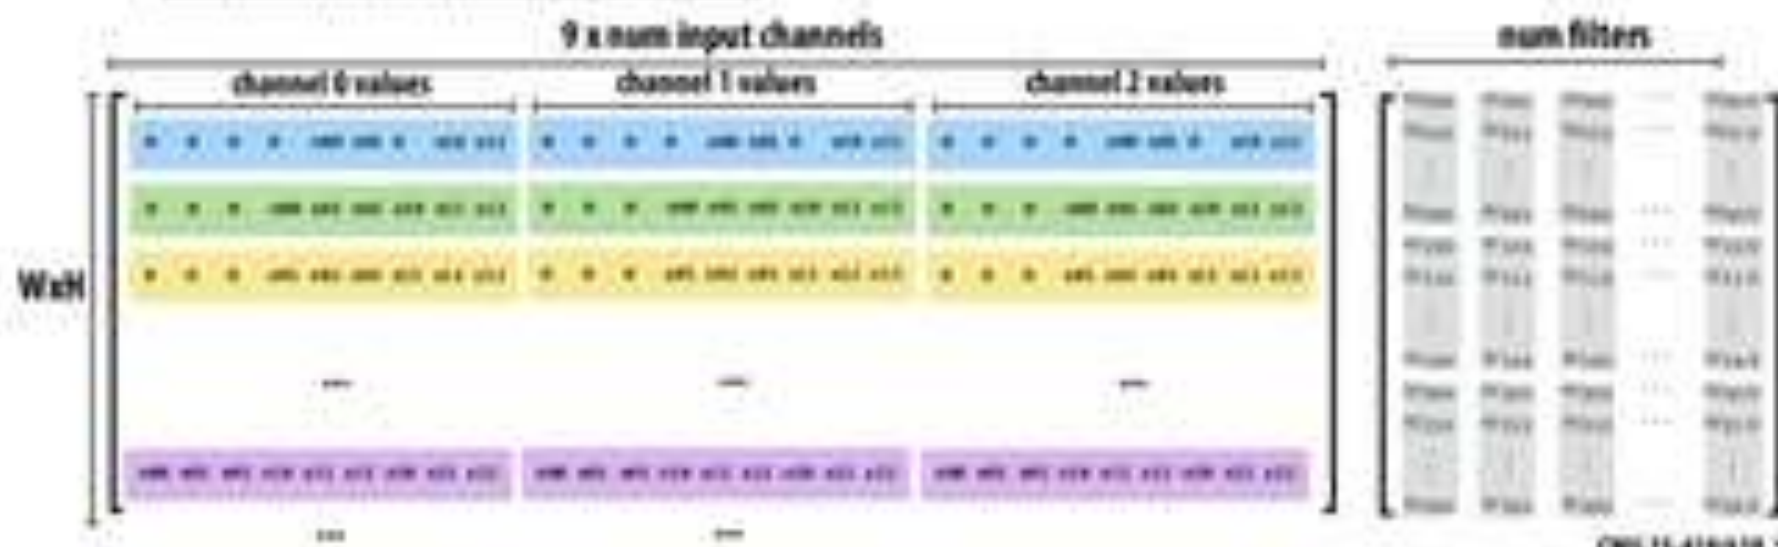
Note: 0-pad matrix

Multiple convolutions on multiple input channels



For each filter, sum responses over input channels

Equivalent to $(3 \times 3 \times \text{num_channels})$ convolution on $(W \times H \times \text{num_channels})$ input data



VGG memory footprint

Calculations assume 32-bit values (image batch size = 1)

inputs/outputs get multiplied by image batch size

multiply by next layer's conv window size to form input matrix to next conv layer!! (for VGG's 3x3 convolutions, this is the data amplification)

	weights mem:	output size (per image)	(mem)
input: 224 x 224 RGB image	—	224x224x3	150K
conv: (3x3x3) x 64	6.5 KB	224x224x64	12.3 MB
conv: (3x3x64) x 64	144 KB	224x224x64	12.3 MB
maxpool	—	112x112x64	3.1 MB
conv: (3x3x64) x 128	228 KB	112x112x128	6.2 MB
conv: (3x3x128) x 128	576 KB	112x112x128	6.2 MB
maxpool	—	56x56x128	1.5 MB
conv: (3x3x128) x 256	1.1 MB	56x56x256	3.1 MB
conv: (3x3x256) x 256	2.3 MB	56x56x256	3.1 MB
conv: (3x3x256) x 256	2.3 MB	56x56x256	3.1 MB
maxpool	—	28x28x256	766 KB
conv: (3x3x256) x 512	4.5 MB	28x28x512	1.5 MB
conv: (3x3x512) x 512	9 MB	28x28x512	1.5 MB
conv: (3x3x512) x 512	9 MB	28x28x512	1.5 MB
maxpool	—	14x14x512	383 KB
conv: (3x3x512) x 512	9 MB	14x14x512	383 KB
conv: (3x3x512) x 512	9 MB	14x14x512	383 KB
conv: (3x3x512) x 512	9 MB	14x14x512	383 KB
maxpool	—	7x7x512	98 KB
fully-connected 4096	392 MB	4096	16 KB
fully-connected 4096	64 MB	4096	16 KB
fully-connected 1000	15.6 MB	1000	4 KB
soft-max		1000	4 KB

Direct implementation of conv layer

```
float input[INPUT_HEIGHT][INPUT_WIDTH][INPUT_DEPTH];
float output[INPUT_HEIGHT][INPUT_WIDTH][LAYER_NUM_FILTERS];
float layer_weights[LAYER_CONV][LAYER_CONV][INPUT_DEPTH];

// assumes convolution stride is 1
for (int img=0; img<IMAGE_BATCH_SIZE; img++)
  for (int j=0; j<INPUT_HEIGHT; j++)
    for (int i=0; i<INPUT_WIDTH; i++)
      for (int f=0; f<LAYER_NUM_FILTERS; f++) {
        output[j][i][f] = 0.f;
        for (int kk=0; kk<INPUT_DEPTH; kk++) // sum over filter responses of input channels
          for (int jj=0; jj<LAYER_CONV; jj++) // spatial convolution (Y)
            for (int ll=0; ll<LAYER_CONV; ll++) // spatial convolution (X)
              output[j][i][f] += layer_weights[f][jj][ll][kk] * input[j+jj][i+ll][kk];
      }
}
```

Seven loops with significant input data reuse: reuse of filter weights (during convolution), and reuse of input values (across different filters)

Avoids $O(N)$ footprint increase by avoiding materializing input matrix

In theory loads $O(N)$ times less data (potentially higher arithmetic intensity... but matrix mult is typically compute-bound)

But must roll your own highly optimized implementation of complicated loop nest.

Conv layer in Halide

```
let in_w, in_h, in_ch = 4;           // input params: assume initialized
func in_func;                         // assume input function is initialized

let num_f, f_w, f_h, pad, stride;    // parameters of the conv layer

func forward = func("conv");
var x("x"), y("y"), z("z"), n("n"); // n is minibatch dimension

// This creates a padded input to avoid checking boundary
// conditions while computing the actual convolution
f_in_bound = BoundaryConditions::repeat_edge(in_func, 0, in_w, 0, in_h);

// Create image buffers for layer parameters
Image(float) w(f_w, f_h, in_ch, num_f)
Image(float) b(num_f);


@doe r(0, f_w, 0, f_h, 0, in_ch);

// Initialize to bias
forward(x, y, z, n) = b(z);
forward(x, y, z, n) += w(r.x, r.y, r.z, z) *
    f_in_bound(x*stride + r.x + pad,
              y*stride + r.y + pad,
              r.z, n);
```

Consider scheduling this seven-dimensional loop nest.

Algorithmic improvements

- Direct convolution can be implemented efficiently in Fourier domain (convolution \rightarrow element-wise multiplication)
 - Overhead: FFT to transform inputs into Fourier domain, inverse FFT to get responses back to spatial domain ($N \log N$)
 - Inverse transform amortized over all input channels (due to summation over inputs)
- Direct convolution using work-efficient Winograd convolutions
1D example: consider producing two outputs of a 3-tap 1D convolution with weights: w_0, w_1, w_2



1D 3-tap total cost:

4 multiplies

8 additions

(4 to compute m 's + 4 to reduce final result)

Direct convolution: 6 multiplies, 4 adds

In 2D can notably reduce multiplications

(3x3 filter: 2.25x fewer multiplies for 2x2 block of output)

$$\begin{bmatrix} y_0 \\ y_1 \end{bmatrix} = \begin{bmatrix} x_0 & x_1 & x_2 \\ x_1 & x_2 & x_3 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} m_1 + m_2 + m_3 \\ m_2 + m_3 + m_4 \end{bmatrix}$$

$$m_1 = (x_0 - x_1)w_0$$

$$m_2 = (x_1 + x_2) \frac{w_0 + w_1 + w_2}{2}$$

$$m_3 = (x_2 - x_1) \frac{w_0 - w_1 + w_2}{2}$$

$$m_4 = (x_1 - x_2)w_2$$

Filter dependent
(can be precomputed)

Reducing network footprint

- Large storage cost for model parameters
 - AlexNet model: ~200 MB
 - VGG-16 model: ~500 MB
 - This doesn't even account for intermediates during evaluation


- Footprint: cumbersome to store, download, etc.
 - 500 MB app downloads make users unhappy!




- Consider energy cost of 1B parameter network
 - Running on input stream at 20 Hz
 - 640 pJ per 32-bit DRAM access
 - $(20 \times 1\text{B} \times 640\text{pJ}) = 12.8\text{W}$ for DRAM access
 - (more than power budget of any modern smartphone)



"Pruning" (sparsifying) a network



"Pruning" (sparsifying) a network



$$f\left(\sum_i x_i w_i + b\right)$$

$$f(x) = \max(0, x)$$

Idea: prune connections with
near zero weight

Remove entire units if all
connections are pruned.

Representing “sparsified” networks

Step 1: prune low-weight links (iteratively retrain network, then prune)

- Over 90% of weights in fully connected layers can be removed without significant loss of accuracy
- Store weight matrices in compressed sparse row (CSR) format

Indices	1	4	9	...
Value	1.8	0.5	2.1	

0	1.8	0	0	0.5	0	0	0	0	1.1	...
---	-----	---	---	-----	---	---	---	---	-----	-----


Reduce storage overhead of indices by delta encoding them to fit in 8 bits

Indices	1	3	6	...
Value	1.8	0.5	2.1	

Efficiently storing the surviving connections

Step 2: Weight sharing: make surviving connections share a small set of weights

- Cluster weights via k-means clustering
- Compress weights by only storing index of assigned cluster ($\lg(k)$ bits)
- This is lossy compression



Step 3: Huffman encode quantized weights and CSR indices (lossless compression)

VGG-16 compression

Substantial savings due to combination of pruning, quantization, Huffman encoding *

Layer	#Weights	Weights% (P)	Weight bits (P+Q)	Weight bits (P+Q+H)	Index bits (P+Q)	Index bits (P+Q+H)	Compress ratio (P+Q)	Compress ratio (P+Q+H)
conv1_1	2K	58%	8	6.8	5	1.7	40.0%	29.91%
conv1_2	37K	22%	8	6.5	5	2.6	9.8%	6.99%
conv2_1	24K	34%	8	5.8	5	2.4	14.3%	8.91%
conv2_2	140K	36%	8	5.9	5	2.3	14.7%	9.31%
conv3_1	255K	53%	8	4.8	5	1.8	21.7%	11.13%
conv3_2	580K	34%	8	4.6	5	2.0	9.7%	5.67%
conv3_3	580K	42%	8	4.6	5	2.2	17.0%	8.96%
conv4_1	1M	32%	8	4.6	5	2.6	13.1%	7.29%
conv4_2	2M	27%	8	4.2	5	2.9	10.9%	5.83%
conv4_3	2M	34%	8	4.4	5	2.5	14.0%	7.47%
conv5_1	2M	35%	8	4.7	5	2.5	14.3%	8.00%
conv5_2	2M	29%	8	4.6	5	2.7	11.7%	6.52%
conv5_3	2M	36%	8	4.6	5	2.3	14.8%	7.79%
fc6	103M	4%	5	3.6	5	3.5	1.6%	1.30%
fc7	17M	4%	5	4	5	4.3	1.5%	1.25%
fc8	4M	23%	5	4	5	3.4	7.1%	5.24%
Total	138M	7.3% (11%)	6.4	4.1	5	3.1	3.2% (11%)	2.05% (48%)

P = connection pruning (prune low weight connections)

Q = quantize surviving weights (using shared weights)

H = Huffman encode

ImageNet Image Classification Performance

	Top-1 Error	Top-5 Error	Model size	
VGG-16 Ref	11.50%	11.52%	552 MB	
VGG-16 Compressed	11.17%	10.91%	11.5 MB	48x

* Benefits of automatic pruning apply mainly to fully connected layers, but many modern networks are dominated by costs of convolutional layers

More efficient topologies

- Original DNNs for image recognition were overprovisioned
- Can think of modern designs as hand-designed to be sparser

SqueezeNet: [Iandola 2017] Reduced number of parameters in AlexNet by 50x, with similar performance on image classification



Inception (GoogleLeNet) — 27 total layers, 7M parameters



ResNet (34 layer version)

Deep neural networks on GPUs

- Many high-performance DNN implementations target GPUs
 - High arithmetic intensity computations (computational characteristics similar to dense matrix-matrix multiplication)
 - Benefit from flop-rich architectures
 - Highly-optimized library of kernels exist for GPUs (cuDNN)
 - Most CPU-based implementations use basic matrix-multiplication-based formulation (good implementations could run faster!)



Facebook's Big Sur

Increasing efficiency through specialization

Example: Google's Tensor Processing Unit (TPU)
Accelerates deep learning operations



Intel Lake Crest ML accelerator
(formerly Nervana)



Emerging architectures for deep learning

- **NVIDIA Pascal (NVIDIA's latest GPU)**
 - Adds double-throughput 16-bit floating point ops
 - This feature is already common on mobile GPUs
- **Intel Xeon Phi (Knights Landing)**
 - Flop rich 68-core x86 processor for scientific computing and machine learning
- **FPGAs, ASICs**
 - Not new: FPGA solutions have been explored for years
 - Significant amount of ongoing industry and academic research
 - Google's TPU

Programming frameworks for deep learning

- Heavyweight processing (low-level kernels) carried out by target-optimized libraries (NVIDIA cuDNN, Intel MKL)
- Popular frameworks use these kernel libraries
 - Caffe, Torch, Theano, TensorFlow, MxNet
- DNN application development involves constructing novel network topologies



nervana
Graph Compiler



Summary: efficiently evaluating deep nets

- Computational structure
 - Convlayers: high arithmetic intensity, significant portion of cost of evaluating a network
 - Similar data access patterns to dense-matrix multiplication (exploiting temporal reuse is key)
 - But straight reduction to matrix-matrix multiplication is often sub-optimal
 - Work-efficient techniques for convolutional layers (FFT-based, Winograd convolutions)
- Large numbers of parameters: significant interest in reducing size of networks for both training and evaluation
- Many ongoing studies of specialized hardware architectures for efficient evaluation
 - Future CPUs/GPUs, ASICs, FPGAs, ...
 - Specialization will be important to achieving "always on" applications